Agency Client Delivery AI Prompt Pack

Agencies: deliver client work 3x faster with 60+ prompt + artifact pairs.

Most AI prompt packs hand you a sentence and call it a win. You paste it in, get a page of generic bullets, and then spend an hour turning that mush into something a client will actually pay for. That hour is where the margin leaks.

This pack is built the other way around. Every prompt in here specifies an **exact deliverable** — the precise artifact structure that lands in the client's inbox: a 4-week content calendar table with named columns, a 12-slide audit report with per-slide content, a campaign brief with 8 defined sections, a monthly report with a fixed executive summary block. You are not asking an AI model for "ideas about content." You are asking it for the finished table, in the right order, with the right cells, ready to paste into Google Docs, Notion, Sheets, or a deck.

What's inside

- **60+ prompts for strategy, content, reporting, proposals and client communication**
- **20 finished artifact templates: content calendars, audit reports, ad copy, campaign briefs**
- **Prompts produce client-facing deliverables, not brainstorming notes**
- **Format: PDF + Notion**

The 68 prompts are grouped into ten categories and numbered so you can cite them internally: *"Strategy prompt 04, run against the discovery transcript."* The 20 artifact templates are the finished structures those prompts output — each one shows the exact sections, columns, row counts, and what belongs in every cell, so you can hand a junior designer or a new account manager a spec instead of a vague brief.

Brand voice is handled at the prompt layer, not bolted on at the end. Any prompt that produces client-visible words accepts `{{brand_voice_guide}}` and is instructed to enforce it — including the enforcement rules (what the voice guide means for sentence length, banned words, jargon, emoji use, and proof style). If the guide is silent on something, the prompt says so rather than inventing a tone.

How to use this pack

The five-step delivery loop

Every deliverable you send a client moves through the same loop. The prompts are designed to sit inside it.

Step	What happens	Prompt support
1. Collect	You gather the raw input: transcripts, analytics exports, brand docs, website copy, past deliverables.	Every prompt opens with a `{{placeholder}}` list so you know exactly what to paste in. If an input is missing, the prompt tells you to flag the gap rather than fabricate.
2. Direct	You set the shape of the output before generation.	Each prompt names the artifact, its sections, row counts, and format. The output format is specified *inside* the prompt so the model can't drift.
3. Generate	You run the prompt and get a structured draft.	Prompts are written for paste-and-run use in ChatGPT, Claude, Gemini, or any similar tool.
4. Shape	You apply your agency's judgment: swap in real numbers, tighten the argument, add screenshots.	The "How to use" line under each prompt tells you what to check and what to replace.
5. Send	The artifact goes to the client as a finished document.	The Client-ready checklist at the end of this pack is the last gate before anything leaves your outbox.

Placeholder reference

Replace every `{{placeholder}}` before you run a prompt. Double braces are the convention throughout this pack. If a placeholder contains a percent sign, a currency symbol, or a client name with an apostrophe, paste it as plain text — the prompts instruct the model to reproduce values exactly as supplied.

Placeholder	What to put in it
`{{client_name}}`	The client's legal or brand name as it should appear on the document.
`{{industry}}`	Industry or vertical — used to calibrate examples, benchmarks, and language.
`{{brand_voice_guide}}`	Your client's voice guide: tone rules, do/don't examples, banned words, reading level. Optional but recommended on every word-producing prompt.
`{{audience}}`	Primary audience description, plus secondary segments.
`{{offer}}`	The product, service, or offer being marketed. Include price point and what's included.
`{{competitors}}`	Named competitors, or the list of brands the client is compared against.
`{{goals}}`	Business goals for the period — revenue, leads, signups, awareness, retention.
`{{kpis}}`	The metrics that define success for this account (e.g. CAC, ROAS, MQL, organic sessions).
`{{metrics_data}}`	Raw analytics pasted in as a table or bullet list. Keep the numbers exactly as exported.
`{{period}}`	The reporting window (e.g. "1–31 May 2026", "Q2 2026").
`{{budget}}`	Media budget, production budget, or project fee.
`{{channel_mix}}`	Planned channel split (e.g. "60% Meta, 25% Google Search, 15% TikTok").
`{{timeline}}`	Working timeline with real dates.
`{{constraints}}`	Legal, regulatory, platform, or brand restrictions.
`{{platform}}`	The specific channel (Instagram, LinkedIn, YouTube, Gmail, etc.).
`{{format_spec}}`	Any house format rules — fonts, slide counts, character limits, file naming.
`{{examples}}`	Past deliverables, past posts, or example copy the client already approved.
`{{website_url}}`	The client's site or a specific page URL.
`{{market}}`	Geographic market, language, and locale (e.g. "UK, English, GBP").
`{{contact_name}}`	The client's stakeholder or recipient.
`{{your_agency}}`	Your agency name — used for report footers, proposal covers, and sign-offs.

Running a prompt

1. Open a fresh chat. Long context dilutes structure — don't paste 40 pages of background and expect a clean table.
2. Paste the prompt code block as-is. The formatting and rules inside it are doing the work.
3. Fill every placeholder. A prompt with `[BRACKETED]` leftovers is a prompt you haven't finished yet.
4. If the output drifts (wrong length, extra sections, missing columns), do not rewrite the whole prompt. Add one line: *"Re-output using the exact structure specified above. Do not add, remove, or rename sections."* That single correction line fixes most drift.

5. Copy the output into your client-facing template — Notion, Google Docs, Sheets, Figma, or Slides — and run the Client-ready checklist.

When the client gives you a mess

Real inputs are messy. A "brief" might be a 600-word voice note transcript. The rule for every prompt in this pack is the same: **the model extracts, labels, and structures what is given, and marks what is missing as `[NEEDS INPUT: ...]` rather than inventing it.** A deliverable with three clearly flagged gaps is faster to finish than one built on six plausible-sounding guesses.

Using the prompts in Notion

Each prompt is designed to be pasted into a Notion AI block, a page property, a database row, or a button that opens a prefilled template. The practical set-up: create one Notion database called `Deliverables` with properties `Client`, `Artifact type`, `Prompt ID`, `Status`, `Due`, and `Link`. The `Prompt ID` field stores the number in this pack (`#14`, `#37`, `#61`) so anyone on the team can find the exact prompt. The `Link` field points at the finished client deliverable — never at the raw AI output. That separation is what stops an unreviewed draft reaching a client.

Working with a team

- **Account managers** own the inputs and the brief.
- **Strategists and creatives** own the artifact itself.
- **QA** owns the Client-ready checklist.
- The `{{brand_voice_guide}}` lives in the client's own Notion page and is pasted fresh for every word-producing prompt — never summarised from memory, because summarised voice guides drift.

A note on models

These prompts are model-agnostic and work in any modern chat assistant. They lean on three behaviours that every capable model supports: strict structural output, explicit placeholder handling, and self-labelled gaps. If your model of choice truncates long tables, ask for the output section by section — "Give me rows 1–15 only" — and stitch them in the client template. Structure survives; prose style varies.

The 20 finished artifact templates

The next section is the heart of the pack. Each template below is a **finished artifact**, not a description of one. What you see is the skeleton your client receives: the exact headings, the exact column headers, the exact number of rows, and a note on what belongs in every cell. When you run the matching prompt, the output should drop into this structure without reshaping.

Two rules apply to all 20:

1. **Cell content is always specific.** "Strong hook" is a note, not a deliverable. "Hook: We audited 240 pages and found 37% of them invisible" is a deliverable.
2. **Every artifact ends with a next action.** A client deliverable that doesn't tell the client what happens next gets no response. Each template's final row or section is the ask.

Template 1 — 30-Day Content Calendar

Artifact format: one table, 30 content rows, plus a header block. Delivered as a Google Sheets–ready table (paste-friendly) or a Notion table. Optionally accompanied by a 4-week month view for layout approval.

Header block (above the table)

Header field	Content
Client	`{{client_name}}`
Period	`{{period}}` — exact start and end dates
Channels	`{{channel_mix}}`
Content pillars	The 3–5 pillars with one-line definitions
Primary goal this month	The single KPI this calendar is built to move
CTA default	The action every post drives toward
Approval path	Who signs off, and by which day each week

Table columns (9 columns, 30 rows — one row per day)

Date	Day	Channel	Pillar	Format	Post type	Hook / first line	Body copy or caption	CTA	Asset + owner
01 May	Thu	Instagram	Education	Carousel	6-slide explainer	"Most teams check GA4 twice a month. Here's why that fails."	Full caption, 90–140 words, opening line repeated, proof point, soft CTA	"Save this for your next review"	`carousel_edu_v1.psd` / Design
02 May	Fri	LinkedIn	Behind the scenes	Text + image	Founder POV	"We were wrong about our best-performing channel."	120–180 words, one idea, one number, one question to close	"What's your most misunderstood channel?"	`linkedin_bts_v1.png` / Social
…									
30 May	Sat	Email	Offer	Newsletter	Monthly roundup	"Everything we shipped in May"	Subject line + preheader + 3 sections + 1 CTA	"See the full breakdown"	`email_may_final.html` / Copy

Fixed rules for this artifact

- Exactly 30 rows. No gaps, no duplicate dates, weekends included.
- No row may have an empty cell. If something is undecided, write `[NEEDS INPUT: ...]` so the gap is visible on the client's side too.
- Each row's hook must be a complete line a human would say out loud, not a topic label.
- Row order is chronological, and the first row's date must equal the first date in `{{period}}`.
- The final row is always a recap/CTA post — no filler.

Template 2 — Weekly Content Calendar

Artifact format: a single-week table (7 rows) designed for fast approval and fast turnaround, plus a two-line summary of the week's angle and goal.

Summary block

Field	Content
Week of	`{{week_starting}}`
Theme / big idea	One sentence the whole week hangs on
Goal	The one outcome this week drives
Reach vs. depth	Whether this week prioritises new audiences or existing engagement
Reused from	Which evergreen asset this week extends (or "new")

Table columns (8 columns, 7 rows — one row per day)

Day	Channel	Post type	Pillar	Hook / first line	Caption or body (final, not outline)	CTA	Asset + due date
Mon	IG	Reel	Product education	"The 30-second setup nobody shows you"	80 words, spoken-style, caption opens with same hook	"Comment 'SETUP'"	`reel_week1_setup.mp4` / Wed
Tue	—	—	—	—	—	—	—
Wed	LinkedIn	Document post	Education	"Our 5-step launch checklist"	150 words + 5 numbered steps	"Repost if useful"	`doc_checklist.pdf` / Tue
Thu	TikTok	Short video	Personality	"Things we got wrong in Q1"	60 words, hook in first 2 seconds	"Follow for the rebuild"	`tt_q1_mistakes.mp4` / Thu
Fri	IG	Story set (4 frames)	Social proof	"Client said this. We kept the receipt."	4 frames, each max 12 words, one sticker CTA	"Read the case study"	`stories_case_study.zip` / Fri

Day	Channel	Post type	Pillar	Hook / first line	Caption or body (final, not outline)	CTA	Asset + due date
Sat	Email	Newsletter	Value	"Three things worth your Sunday"	Subject, preheader, 3 items, 1 CTA	"Reply with your weekend win"	`email_wk3_sunday.html` / Sat
Sun	—	Rest / batch	—	"No posting — batch Monday's assets"	n/a	n/a	—

Fixed rules for this artifact

- Exactly 7 rows. A blank day uses an em dash in every cell and a stated reason in the Post type cell (rest, batch, campaign hold) — never a silent blank.
- Captions are written in final copy, ready to paste. No bullet outlines of what to say.
- The table is built to be approved on a phone screen, so each cell stays under two lines.
- If the week contains a launch, the launch row is marked with a 📌 in the Day cell.

Template 3 — SEO Audit Report

Artifact format: a structured report, 12 sections, with a 14-slide deck equivalent. Designed to be read in 20 minutes and acted on without a call.

Report sections

#	Section	Contains
1	Cover	Client, site URL, audit date, audit scope, preparer name
2	Executive summary	5 bullets max — traffic trend, biggest loss, 3 priority fixes, 1 risk
3	Site health snapshot	Traffic, indexed pages, average position, clicks, impressions, Core Web Vitals pass rate
4	Technical findings	Table of issues with severity, affected URLs, fix, owner, effort
5	On-page findings	Per-page title tag, meta description, H1, word count, internal links, status
6	Content gap analysis	Keywords with high volume and thin coverage, mapped to existing pages
7	Keyword performance	Current ranking positions for {{target_keywords}}, with movement vs. prior period
8	Backlink profile	Referring domain count, new/lost links, top 5 links, toxic signals
9	Competitor comparison	Side-by-side of 3 competitors on the same metrics
10	Structured data	Detected schema, gaps, errors, opportunity
11	Prioritised roadmap	90-day plan: quick wins, mid-term, structural
12	Measurement plan	KPIs, tools, reporting cadence, and the first 3 checkpoints

Core findings table (used in sections 4, 5, 6 — same column set throughout)

#	Issue	Severity	Pages affected	Evidence	Recommended fix	Owner	Effort	Impact score
1	Missing meta descriptions	High	47	Screaming Frog export	Write unique descriptions from intent + {{brand_voice_guide}}	Copy	M	8/10

Fixed rules for this artifact

- Severity is one of `High / Medium / Low` and is defined in the report header.
- Impact score combines traffic upside and effort; the methodology is stated in one line.
- Every recommendation names an owner and an effort estimate. Recommendations without owners do not ship.
- Numbers are dated and sourced: "GA4, 1–31 May 2026" or "Ahrefs, pulled 4 June 2026".

Template 4 — Social Media Audit

Artifact format: a channel-by-channel audit across 4 platforms, with a scorecard, a top-issues table, and a 30-day corrective plan.

Scorecard (one row per channel + one roll-up row)

Channel	Followers	Posting cadence (actual vs. planned)	Avg. engagement rate	Reach (30d)	Saves	Shares	Follower growth	Best-performing post (why)	Score /10
Instagram	4,210	8 vs 12 planned	3.1%	86,400	310	92	+84	Carousel on "3 ways to..." — saves drove reach	6/10
LinkedIn	1,980	3 vs 4 planned	2.4%	22,100	180	44	+31	Text post naming a mistake — comments from peers	7/10
TikTok	640	1 vs 5 planned	5.8%	41,900	620	210	+52	Hook in first 1.5s, native audio	5/10
X	300	0 vs 3 planned	0.0%	0	0	0	-4	Channel abandoned mid-pilot	1/10
Roll-up									**Needs a cadence decision, not more content**

Top issues table (7 rows)

#	Issue	Evidence	Why it hurts	Fix	Effort	Owner	Deadline
1	Cadence gap on TikTok	4 of 5 planned videos unshipped	Disrupts the algorithm's learning	Batch 2 videos/week on Mondays	Low	Social lead	12 Jun

30-day corrective plan

Week	Theme	Deliverables	Channel focus	Target metric
Week 1	Fix the pipeline	3 videos, 8 posts, 1 email	TikTok + Instagram	Publishing cadence met 5/5 days
Week 2	Fix the hook	3 videos, 8 posts, 1 email	Instagram + LinkedIn	+25% saves
Week 3	Fix the CTA	3 videos, 8 posts, 1 email	All	+0.4pp engagement rate
Week 4	Fix the proof	3 videos, 8 posts, 1 email	LinkedIn + email	Email click rate ≥ 3%

Fixed rules for this artifact

- Score out of 10 with the scoring criteria stated (reach, consistency, engagement quality, conversion path, brand consistency).
- Every "best-performing post" is explained with a *why*, not just ranked.
- The audit never recommends "post more." It recommends a specific fix for a specific constraint.
- Ends with a one-paragraph recommendation and a single question for the client.

Template 5 — Ad Creative Matrix (Meta + Google)

Artifact format: two matrices in one artifact — a Meta matrix and a Google matrix — each cell containing finished ad copy, not a copy angle.

Meta matrix (5 columns, 4 rows of angles × 3 formats = 12 ads)

Angle \ Format	Reel / Video (9:16)	Carousel	Static / Image
Problem-agitate	Hook: "Your checkout loses 71% of mobile carts." Body: 65 words. CTA: "See the fix"	6 cards: pain → 3 causes → 1 stat → fix → proof → CTA	Headline: "Mobile carts, fixed." Copy: 40 words. Proof badge + logo lockup
Proof / social	30s testimonial cut, caption overlay, first frame = outcome	5 cards: quote → metric → how → result → CTA	Testimonial card, brand palette, quote as headline
Comparison	Split-screen before/after, 20s	5 cards: them vs us across 4 rows → CTA	Two-column comparison table graphic
Offer / urgency	15s product demo, price on screen, countdown	4 cards: value → price → deadline → CTA	Price + deadline + guarantee badge

Each matrix cell must also record: primary text (final copy), headline, description, CTA button, format spec (dimensions, file size), thumbnail direction, and the audience/intent it serves.

Google matrix (Search + Performance Max + YouTube)

Ad type	Headlines (15 max, 30 chars each)	Descriptions (4, 90 chars each)	Path fields	Asset group	Notes
Responsive Search	15 final headlines	4 final descriptions	2 paths, 15 chars each	Group A — high intent	Pin positions 1 and 2
Performance Max	10 headlines, 4 long headlines, 4 descriptions, 5 image prompts	—	—	Group B — prospecting	Add audience signals
YouTube	—	2 descriptions	—	Group C — video	15s + 30s cuts

Required companion section: an **Asset production sheet** with one row per asset (17 rows for the full set), columns: `Asset ID | Platform | Format | Spec | Copy (final) | Design brief | Due | Status`, and a **Testing plan** table: `Test | Hypothesis | Variable | Metric | Decision threshold | Decision date`.

Fixed rules for this artifact

- Every headline respects the platform character limit, with the character count written next to it: `Free shipping over $50 (29)`.
- No angle appears twice with the same copy. Twelve Meta ads means twelve genuinely different executions.
- Every ad maps to a funnel stage and a target audience, both named.
- The testing plan names a decision threshold in advance. "See how it does" is not a threshold.

Template 6 — Campaign Brief

Artifact format: an 8-section brief document, 2–3 pages, that locks the campaign before any asset is made. Signed off by both sides.

#	Section	Contains	Typical length
1	Campaign snapshot	Name, one-line proposition, dates, budget, markets, channels, owner	One table + 3 lines
2	Background & context	What the business is doing now, what this campaign must change, why now	150–250 words
3	Audience	Primary segment, secondary segment, insight, tension, what they believe today, what they need to believe after	1 table + 200 words
4	Objectives & KPIs	Business objective → campaign objective → KPI → target → measurement source	1 table, 3–5 rows
5	Strategy	The single-minded proposition, the strategic idea, the 3 supporting proof points, the message hierarchy	1 boxed statement + list
6	Creative direction	Tone, visual world, do/don't, mandatory elements, reference territory, what we're avoiding	1 table + 150 words
7	Channel plan	Per channel: role in funnel, format, volume, cadence, asset list, budget split	1 table, one row per channel
8	Plan, budget & next steps	Phased timeline with milestones, budget allocation table, RACI, approval gates, 5 next actions	2 tables + list

Section 1 table

Field	Value
Campaign name	{{campaign_name}}
Proposition	The one sentence the whole campaign says
Flight dates	{{start_date}} → {{end_date}}
Total budget	{{budget}} + channel split
Markets / language	{{market}}
Channels	{{channel_mix}}
Campaign owner	{{your_agency}} + client-side approver

Section 4 table

Business objective	Campaign objective	KPI	Baseline	Target	Source
Grow qualified pipeline	Generate demo requests from qualified accounts	MQLs from paid	40/mo	75/mo	HubSpot, GA4

Section 7 table

Channel	Funnel role	Format	Volume	Cadence	Assets required	Budget	Success metric

Section 8 timeline

Phase	Dates	Milestones	Approval gate	Owner

Fixed rules for this artifact

- Section 5 contains exactly one single-minded proposition. Two ideas means no idea.
- Every creative instruction is either a "do" or a "don't." No adjectives without a reference.
- Section 8 ends with a numbered list of the next 5 actions, each with an owner and a date.
- No section is left as "TBC" — unfilled decisions are written as `[DECISION NEEDED: ...]` with the decision-maker named.

Template 7 — Client Strategy One-Pager

Artifact format: a single page, front and back maximum, designed to be read in 90 seconds before a meeting. Two sides, seven blocks.

Side A

Block	Contains	Word limit
Header	Client, period, prepared by, date	—
Situation	3 lines on what's true now — market, business, channels	60 words
Insight	The one non-obvious thing we believe, and the evidence behind it	60 words
Objectives	3 objectives, each with a numeric target	3 bullets

Side B

Block	Contains	Word limit
Strategy statement	One sentence: how we win, expressed as a choice	35 words
Strategic pillars	3 pillars — name, what it means, what we'll do, what success looks like	3 × 45 words
Audience priority	Who we're prioritising and who we're deprioritising, with a reason	60 words
90-day priorities	The top 3 moves with owner and first milestone	3 bullets

Fixed rules for this artifact

- Fits on one printed page — if it doesn't, the content is wrong, not the font size.
- Every claim is traceable to a source the client saw (data, research, or an agreed assumption).
- The deprioritisation is stated. A one-pager that only says what you'll do isn't a strategy.
- Exactly one insight. If there are two, one is a tactic.

Template 8 — Monthly Performance Report

Artifact format: the report every client expects to receive on a set day each month. 10 sections, with a 1-page executive summary that can be forwarded on its own.

#	Section	Contents
1	Cover	Client, period, prepared by, date, reporting cadence
2	Executive summary	150 words max: the month's headline, the single best result, the single biggest miss, what we're changing
3	KPI scorecard	All KPIs vs. target vs. prior period vs. trend arrow
4	Channel performance	Per-channel table with spend, reach, clicks, leads, CPA/ROAS, efficiency rating
5	Top performing creative	3 assets with the numbers, why they worked, and what we learned
6	Underperformers & losses	3 assets/campaigns with the numbers, what we think happened, and the decision (pause/iterate/kill)
7	Creative production log	Everything shipped this month, with dates and performance
8	Insights & observations	3–5 patterns found in the data, each with supporting numbers
9	Next month's plan	Actions, channels, budget, campaign ideas, dates
10	Action required	3–5 items where the client needs to decide, approve, or provide something

Section 3 scorecard

KPI	Target	Actual	vs target	vs last month	Trend	Status
Qualified leads	75	82	+9%	+14%	▲	On track
CAC	≤ $120	$138	+15%	+3%	▲	At risk

Section 4 channel table

Channel	Spend	Impressions	Reach	Clicks	CTR	Conversions	CPA	ROAS	Efficiency rating

Section 10 action table

#	Action needed	Why it matters	Decision owner	Needed by	Impact if delayed

Fixed rules for this artifact

- Every number carries its source and date. If a number is an estimate, it says so.
- Prior-period comparison is mandatory; a client cannot judge a month in isolation.
- Section 6 always states a decision. A report that identifies losers without deciding their fate will be repeated next month.
- Section 2 is written to be forwarded to the client's boss without editing — no jargon about the campaign, just outcomes.
- If a metric is unavailable, write `Not available — reason` rather than omitting the row.

Template 9 — Quarterly Business Review Deck Outline

Artifact format: a 14-slide deck outline. Each slide lists the headline message (the slide title is the takeaway, not a label), the content blocks, the visual, and the speaker note.

Slide	Headline message (example)	Content blocks	Visual	Speaker note focus
1	Q2 grew pipeline 38% — here's what drove it	Title, client, date, prepared by	Cover	One-sentence framing of the quarter
2	Three things to know before we go further	Agenda: results, what worked, what's next, decisions	Numbered list	Keep to 60 seconds
3	Revenue: 38% above plan, concentrated in two segments	Actual vs target chart, revenue table, segment breakdown	Bar chart + table	Explain the concentration
4	Acquisition cost fell 22% as creative volume doubled	CAC trend line, spend vs CAC, creative volume	Dual-axis chart	Connect spend to efficiency
5	Paid social contributed 61% of pipeline	Pipeline contribution by channel	Stacked bar	Name the assumption
6	Search is underperforming on head terms	Rank movement, CTR, impression share	Line + table	Set up the SEO section
7	Three creative concepts drove 44% of conversions	Top 3 assets with spend and CVR	Asset cards	Tell the story behind one
8	What we tested and what we learned	Test table: hypothesis, result, decision	Table	Note the failed tests too
9	Where we lost money	Waste table: asset, spend lost, cause, action	Table	Take ownership
10	The market shifted: competitor X launched in May	Competitive movement, share, positioning	Comparison table	So what for us?
11	Q3 plan: three priorities, one constraint	Objectives, pillars, budget	Three-column	Name the constraint honestly

Slide	Headline message (example)	Content blocks	Visual	Speaker note focus
12	Q3 roadmap	Month-by-month plan	Gantt-style	Highlight two deadlines
13	Four decisions we need from you	Decision requests with options and recommendations	Table	Each has a recommendation
14	Appendix	Method, data sources, definitions	Index	Reference only

Decision-request table (slide 13)

#	Decision	Options	Our recommendation	Needed by	Consequence of no decision

Fixed rules for this artifact

- 14 slides maximum. If the content doesn't fit, it belongs in the appendix.
- Slide titles are sentences that could be pasted into an email — no labels like "Results."
- Every slide has a speaker note. A deck read without context loses its nuance.
- The deck opens with results, not with the agenda's value. Slide 2 may be skipped entirely for a 20-minute version.
- Numbers on slides are rounded for presentation; the appendix carries the precise figures.

Template 10 — Website Copy Page

Artifact format: a single page's worth of finished copy, structured section by section, ready to paste into the CMS. Includes SEO metadata and an accessibility note per section.

Metadata block

Field	Content	Spec
URL slug	/features/reporting	Lowercase, hyphenated, max 5 words
Page title	Reporting that answers the question	≤ 60 chars
Meta description		≤ 155 chars, includes primary keyword once
Primary keyword		From {{keyword_research}}
Search intent	Informational / commercial / transactional / navigational	—
Internal links (in)		3–5, with anchor text
Internal links (out)		2–4
CTA	Primary + secondary	

Page copy sections

Section	Element	Copy (final)	Notes
Hero	H1		One line, states the outcome, not the feature
Hero	Subhead		20–30 words, expands the promise, includes the mechanism
Hero	Primary CTA		Verb-first, 2–4 words
Hero	Trust line		Logos, rating, or guarantee

Section	Element	Copy (final)	Notes
Problem	H2		Frame the cost of the status quo
Problem	Body		60–90 words, use the client's own language from discovery
Problem	Stat callout		One credible number with source
Solution	H2		Position
Solution	Body		80–120 words, how it works
Solution	3 benefits		Feature + outcome + proof, one line each
How it works	H2 + 3 steps		Step name, one sentence, one detail
Social proof	H2		
Social proof	Quote		Verbatim, with name, role, company
Social proof	Result stat		Number + baseline + timeframe
Objection handling	H2 + 3 objections		The question in the client's voice, then the answer
FAQ	5 questions		Question + 40-word answer
Final CTA	H2 + body + button		Repeat the promise, add the reason to act now

Fixed rules for this artifact

- Copy is final, not outline. If the section needs a placeholder for a real number, it uses `[CLIENT STAT: ...]`.
- The hero H1 is under 12 words and contains no adjective that isn't doing work.
- Every benefit line is `feature → outcome`, never just the feature.
- Objection handling uses the exact words clients used in the discovery call, not the agency's phrasing.
- CTA language is consistent across the whole page — the same button text everywhere.

Template 11 — Landing Page Wireframe Spec

Artifact format: a build-ready spec: every section, its purpose, layout, content inventory, and interaction notes — enough for a designer to lay out and a developer to build without a call.

Spec table — one row per section (12 rows typical)

#	Section	Purpose	Layout	Grid	Content inventory	Interactions	Assets needed	Priority
1	Announcement bar	Drive traffic to a campaign	Full-width, 40px	12 col	1 line + link	Click-through to {{campaign_url}}	None	P1
2	Hero	Explain offer + capture action	Split 6/6	12 col	H1, subhead, CTA, form, trust logos	Form validation, button states	Hero image, 3 logos	P1
3	Logo strip	Borrow credibility	5 logos, single row	12 col	5 SVG logos	Hover state	5 files	P2
4	Problem	Create recognition	Text block + illustration	7/5	H2, 90 words, 3 bullets	None	Illustration	P1

#	Section	Purpose	Layout	Grid	Content inventory	Interactions	Assets needed	Priority
5	Value props	Communicate benefits	3 cards	4/4/4	3 × (icon, title, 30 words)	Hover lift	3 icons	P1
6	Product visual	Show the product	Full-bleed image + 3 callouts	12 col	Screenshot, 3 annotation labels	Scroll-linked callouts	1 screenshot + 3 PNGs	P1
7	Social proof	Reduce risk	2-column quote cards	6/6	2 quotes + 2 logos	None	2 portraits	P1
8	Comparison	Clarify choice	Table, 4 rows	8/4	Us vs alternative × 4	None	Icons	P2
9	Pricing	Remove price friction	2 tiers	6/6	2 tier cards + features	Toggle monthly/ annual	2 graphics	P1
10	FAQ	Handle objections	Accordion, 5 items	8/4	5 Q+A	Expand/collapse	None	P2
11	Lead form	Capture	Sticky bottom bar + inline	12 col	4 fields + consent	Field validation, error states	None	P1
12	Final CTA	Convert	Centred, high contrast	12 col	H2, 1 line, button	—	Background image	P1

Per-section detail block (repeated for each of the 12 rows)

Field	Content
Section name	Exact name used in Figma layer naming and in the ticket
Purpose	One line — what this section does for the visitor
Layout	Structure description with breakpoints (desktop 1440 / tablet 768 / mobile 375)
Content inventory	Every copy block and asset, with source (which artifact it comes from)
Interaction	Behaviour and state changes
Accessibility	Contrast note, alt text, keyboard behaviour, heading level
Copy source	Which prompt/artifact produced this copy (e.g. Prompt 44 → landing page headline set)

Fixed rules for this artifact

- Every section states its purpose in one line. Sections without a purpose get cut.
- Priorities (P1/P2/P3) let the client understand what ships first.
- Accessibility is not optional — every spec includes heading level, contrast and alt text.
- Copy source is recorded so any line can be traced back and edited in one place.

Template 12 — Email Newsletter

Artifact format: one complete email, from subject line to sign-off, plus the sending metadata and a QA strip. Designed for a client to approve and hit send on.

Sending metadata block

Field	Content
Send date / time	{{send_date}}, {{send_time}} + timezone
Audience segment	{{segment}}
Send type	Broadcast / triggered
From	Name + reply-to
Subject line	Final text
Preheader	Final text, 40–90 chars
Primary CTA	
Tracking	UTM set, link tracking on/off, open pixel, click map

Email body structure

Block	Element	Copy (final)	Length
1	Eyebrow / label		2–4 words
2	Headline		Under 12 words
3	Intro		40–60 words, personal and specific
4	Feature / story 1	Heading + 60 words + CTA	
5	Feature / story 2	Heading + 60 words + CTA	
6	Feature / story 3	Heading + 60 words + CTA	
7	Proof	Quote or stat	
8	CTA block	Button + supporting line	
9	Sign-off	Name, role, org, contact	
10	Footer	Unsubscribe, address, preferences	

Copy block (repeated per story)

Field	Content
Heading	
Body	60 words, one idea, one link
Link text	Descriptive, not "click here"
Image / thumbnail	Description + alt text
CTA	

QA strip (bottom of the document, not of the email)

Check	Status
Subject line under 45 chars	✓
Preheader does not duplicate subject	✓
All links tested and UTM-tagged	✓
Plain-text version included	✓
Alt text on all images	✓
Unsubscribe link present	✓
Mobile render checked (375px)	✓
From name recognised by recipients	✓
Link text meaningful out of context	✓

Fixed rules for this artifact

- Subject line includes one specific, concrete detail — a number, a name, a date, or a result.
- Link text describes the destination. "Read more" is not acceptable.
- The email is written to be read on a phone: short paragraphs, one idea per block.
- The sign-off uses a real person's name and role, never "The Team" unless it genuinely is one.

Template 13 — Brand Voice Guide

Artifact format: a working voice guide, 5 sections plus a reference appendix, that anyone can write with. Not a personality description — a set of rules with examples.

#	Section	Contains
1	Voice in one line	A single sentence describing how the brand sounds, with a named reference (person, brand, or vibe)
2	The 4–6 voice traits	Each trait: definition, sounds like / doesn't sound like, before-after example
3	Writing rules	8–12 concrete rules with a good/bad pair for each
4	Words we use / don't use	Two columns, 15–25 entries each, with the reason for every "don't"
5	Mechanics	Reading level, sentence length, paragraph length, punctuation, capitalisation, emoji, numbers, dates, acronyms
Appendix	Worked examples	3 rewrites of the same message in the brand voice, plus 2 examples of what not to do

Section 2 trait table

Trait	What it means	Sounds like	Doesn't sound like	Before → After
Plainly expert	We know the topic and explain it without hiding behind jargon	"Here's what changed and what to do about it."	"Leverage our best-in-class solutions."	"We fixed the checkout bug." → "Your checkout now works on mobile."
Calmly direct	We say the thing, once, without hype	"This didn't work. Here's why."	"We're absolutely thrilled to announce!!"	"This didn't work. Here's why."

Section 3 writing-rules table

#	Rule	Do	Don't
1	One idea per sentence	"We rewrote the guide."	"We rewrote, restructured, and relaunched the guide."
2	Active voice	"The team shipped it."	"It was shipped by the team."

Section 4 word table

We use	We don't use	Why
help	facilitate	"Help" is what people actually say
use	leverage	Jargon that hides the action
broken	suboptimal	Say the actual state

Appendix worked-example table

#	Source text (off-voice)	Rewrite (in voice)	Which rules applied
1	"We're excited to announce our innovative new analytics suite!"	"You can now see where every lead comes from."	1, 4, 6

Fixed rules for this artifact

- Every trait has a before/after example. A trait without an example is a slogan.
- The "we don't use" column states the reason — "avoid it" alone will be ignored.
- Mechanics include a numeric target (reading level, sentence length) so QA can check objectively.
- The guide is 2–4 pages. Longer voice guides don't get used.

Template 14 — Persona Sheet

Artifact format: one persona per sheet, with a fixed 12-block layout. Create one sheet per segment; keep them comparable.

#	Block	Content	Notes
1	Persona name	A first name + role	Not "User A" — a name makes the persona usable
2	One-line snapshot	"Who they are, in one sentence."	15–20 words
3	Role & context	Job title, company size, industry, tenure	Fact-based
4	Day in the life	A real description of their week, the tools they use, who they report to	80–120 words
5	Goals	3 job goals, ranked	"What they're paid to achieve"
6	Challenges	3–4 day-to-day problems	
7	Pain points	3 specific, emotionally real frustrations	
8	Desired outcome	What "good" looks like for them, in their words	

#	Block	Content	Notes
9	Buying triggers	The 4 events that make them start looking	
10	Objections	3 real objections with the response to each	
11	Content habits	Where they spend time, what they follow, what they ignore	
12	Messaging implications	3 direct "so what" statements for the marketing team	This is the operative block

Section 12 table

Message	Evidence it comes from	Where to use it
"Show me the before and after, not the setup."	Pain point 2	Landing page proof section

Fixed rules for this artifact

- Every persona is built from real evidence: call transcripts, survey answers, support tickets, sales notes, or CRM fields. The evidence column is not optional.
- Goals and challenges stay distinct: a goal is what they're trying to achieve, a challenge is what's in the way.
- Section 12 always contains at least one line that constrains what the team should make. A persona with no constraints is a brochure.
- Personas are never "the audience." They are named people with a context.

Template 15 — Competitive Analysis

Artifact format: a landscape comparison across 5 competitors, followed by a positioning conclusion. The conclusion is what the client is paying for.

Comparison matrix (5 competitors + the client = 6 rows, 10 columns)

Brand	Positioning line	Primary audience	Offer / price point	Key promise	Top 3 channels	Content cadence	Strongest asset	Weakest gap	Proof used
{{client_name}}									
Competitor A									

Per-competitor deep dive (one block per competitor, 5 blocks)

Field	Content
Who they are	One line, category and positioning
What they do well	3 points, each with evidence
Where they fall short	3 points, each with evidence
Their signature move	The tactic or message they own
Their most recent change	What launched, when, and the visible effect
Who's switching to them	The segment most likely to leave the client for them

Positioning conclusion

Field	Content
Category we're claiming	
Our differentiator	One line
Proof that supports it	3 items
Message we should not make	What we can't credibly claim
Competitive response	What competitors will do in response to our push, and our counter

Evidence table (used in deep dives)

Claim	Evidence	Source	Date checked	Confidence

Fixed rules for this artifact

- Every competitive claim has evidence and a date checked. Prices, offers, and messaging change — undated claims rot.
- The matrix compares on positioning and proof, not on feature checkboxes nobody sells.
- The conclusion states what we should *not* claim — the most valuable line in most of these documents.
- The client appears in the matrix too, assessed as honestly as the competitors.

Template 16 — Keyword Research Map

Artifact format: a map that connects search demand to the pages that should exist — not a flat list of keywords with volumes.

Map structure

Stage	Rows	Columns
1. Demand clusters	8–12 clusters	Cluster, intent, total monthly volume, difficulty, business relevance, priority
2. Keyword-to-page map	30–50 keywords	Keyword, cluster, intent, volume, difficulty, current ranking URL, target URL, action
3. Content gap	10–15 gaps	Topic, keywords included, why it's missing, recommended new/existing page, priority
4. Cannibalisation check	Any conflicting set	Two URLs ranking for the same intent, fix recommendation
5. Quick wins	8–12	Keyword, position (4–20), volume, difficulty, action, expected click gain

Stage 1 table

Cluster	Primary intent	Monthly volume	Avg. difficulty	Relevance (1–5)	Priority (A/B/C)

Stage 2 table

Keyword	Cluster	Intent	Volume	KD	Currently ranking	Target URL	Action

Stage 3 table

Missing topic	Keywords in cluster	Why it's missing	Recommended page	New or expand	Priority

Stage 4 table

Intent	URL A	URL B	Conflict type	Fix

Stage 5 quick-wins table

Keyword	Position	Volume	KD	Action	Expected clicks/mo

Fixed rules for this artifact

- Every keyword maps to a URL — existing or proposed. A list with no destination is a list, not a map.
- Intent is assigned before priority; intent determines the page type.
- Volume and difficulty figures carry their tool and date.
- Quick wins are chosen on position, not volume alone: a #6 keyword with 800 volume beats a #40 keyword with 20,000.

Template 17 — Influencer Brief

Artifact format: a creator-ready brief that gets opened and followed, plus a content-rights and reporting spec.

Brief document

Section	Contents
Header	Campaign, creator name, dates, contact, deliverables count
Why you	2 sentences on why this creator was chosen, referencing their actual content
The one idea	The single thing the content must land
Your role	What the creator is: a voice, a demonstrator, a reviewer — pick one
Deliverables	Table: format, platform, quantity, length, due date, usage rights
Talking points	3–5, phrased as things to say, not topics
Do / don't	5 rows each, creator-specific
Required elements	Product name, link, disclosure, hashtag, date mention — with exact wording
CTA	Exact button text and link destination
Claim boundaries	What must not be said (claims, comparisons, guarantees)
Music & visuals	Reference posts, what to avoid
Approval process	Rounds, timeline, who approves, turnaround expectation
Payment & terms	Fee, bonus structure, usage rights duration, exclusivity window

Deliverables table

#	Format	Platform	Quantity	Length	Due	Usage rights	Paid

Do / don't table

	Do	Don't
1		

Required-elements table

Element	Exact wording / placement
Disclosure	"#ad — paid partnership with {{client_name}}"
CTA	"Use code {{code}} at checkout"

Fixed rules for this artifact

- The "why you" section references specific content the creator has made. Generic praise signals a templated brief and creators ignore it.
- Deliverables are specified in quantity, length, platform, and date. "One video" is not a spec.
- Claim boundaries are explicit, especially for regulated categories.
- Usage rights are stated in the brief, not negotiated after the content is live.

Template 18 — Creative Brief

Artifact format: a two-page brief that briefs a designer or an AI image/video tool and removes all interpretation. 9 sections.

#	Section	Contains
1	Assignment	One line: what we're making, for whom, for where
2	Background	60–90 words of context the creator needs but shouldn't have to research
3	Audience	Who it's for, what they know, what they're sceptical of
4	Objective & single-minded proposition	The one thing the piece must communicate
5	Message hierarchy	3 tiers: primary message, supporting points (3–4), proof points
6	Tone & voice	Reference, plus the voice rules that apply here
7	Visual direction	Reference territory, mandatory elements, colour, typography, composition, imagery style
8	Mandatory copy	Every word that appears on the asset, final, in order
9	Deliverables, specs & timing	Formats, dimensions, file names, deadlines, approval rounds

Section 7 visual-direction table

Element	Direction	Reference	Mandatory?
Colour	Warm neutral base, single accent for CTA	{{brand_colours}}	Yes

Element	Direction	Reference	Mandatory?
Type	One family, 2 weights, no italics	—	Yes
Subject	Real hands, real environment, no stock smiles	Reference post 1	Yes
CTA treatment	Accent button, 44px tall minimum	—	Yes

Section 8 mandatory-copy table

Order	Element	Exact copy	Character limit	Placement
1	Headline			Top third
2	Subhead			Below headline
3	CTA			Bottom right

Section 9 deliverables table

#	Format	Dimensions	Aspect	File name	Due	Variants

Fixed rules for this artifact

- Section 8 contains final copy only. Placeholder text in a creative brief ships as placeholder text.
- Every visual instruction is either mandatory or optional — nothing sits in between.
- File naming is specified, because 90% of "missing asset" messages are naming problems.
- Reference territory is links or named existing assets, never "make it premium."

Template 19 — Creative Concept Board

Artifact format: a board that presents 3 concepts side by side, each with a stated rationale, so a client chooses rather than reacts. 3 concepts × 10 blocks.

Board overview (one row per concept)

#	Concept name	One-line idea	Format	The hook	Why it fits the brief	Production effort	Risk	Recommendation
1	"The Audit"							
2	"The Switch"							
3	"The Receipt"							

Per-concept block (repeated 3 times)

Block	Contents
Concept name	Memorable, 2–3 words
The idea	2 sentences explaining the mechanic
Why this audience	The specific human truth it taps
Visual storyboard	4 frames, described in one line each

Block	Contents
Sample copy	3 lines of real copy as it would appear
Why it works	3 bullets tied to the brief's objective
Risk / what could go wrong	1–2 lines
Production notes	What it needs, cost, time, assets
Success signal	The metric this concept is designed to move
Deliverables required	Formats and variants

Storyboard table (per concept)

Frame	What's on screen	Text overlay	Purpose
1			Hook
2			Context
3			Proof
4			CTA

Fixed rules for this artifact

- Exactly 3 concepts. More than 3 causes choice paralysis; fewer gives no comparison.
- Each concept includes its risk. A board without risks reads as a sales pitch.
- One concept is recommended, with the recommendation stated in the overview row.
- Sample copy is written, not described — the client should be choosing from finished words.

Template 20 — Client Onboarding Plan

Artifact format: a 4-week onboarding plan that runs from kick-off to a signed-off first plan, with a shared tracker, access checklist, and risk list.

Plan overview

Week	Milestone	Owner	Client dependency	Deliverable at end of week
Week 1	Access & discovery complete	Agency + Client	Ad accounts, analytics, brand assets	Discovery summary + access tracker green
Week 2	Strategy & positioning agreed	Agency	Stakeholder time	Strategy one-pager signed off
Week 3	Creative concepts approved	Agency + Client	Review time from client	Concept board + approved routes
Week 4	First calendar & launch plan locked	Agency	Final sign-off	30-day calendar + launch checklist

Access checklist

#	Asset / account	Owner	Needed by	Status	What breaks without it
1	Meta Business Manager access	Client	Day 2	Not started	No paid social launch
2	GA4 + GTM access	Client	Day 2	Not started	No baseline, no reporting
3	Brand assets (logo, fonts, colours)	Client	Day 4	Not started	Design work blocked
4	Previous 12-month performance export	Client	Day 4	Not started	No historical context

Week-by-week detail (repeated for 4 weeks)

Field	Content
Theme	The purpose of the week
Agency tasks	4–6 specific tasks
Client tasks	2–4 specific tasks, with names if known
Meeting	Agenda with 5 items, attendees, duration
Deliverable	The artifact produced, and the artifact template it matches
Exit criteria	The condition that must be true to close the week
Risk	The most likely thing to go wrong, and the mitigation

Kick-off agenda

#	Item	Minutes	Owner	Output
1	Introductions and roles	5	Both	Named contacts list
2	Business goals and KPIs	15	Client	Agreed goals list
3	What has and hasn't worked	15	Client	Prior learnings doc
4	Scope and what's out of scope	10	Agency	Agreed scope statement
5	Access, timelines, and comms	10	Both	Access checklist + cadence

Standing cadence table

Meeting	Frequency	Duration	Attendees	Agenda source	Who sends notes

Fixed rules for this artifact

- Every week ends with a deliverable and an exit condition. "Relationship building" is not a deliverable.
- Client tasks are named and dated, so delay is visible early rather than at week 4.
- The kick-off agenda has a time budget per item; total equals the meeting length.
- A standing risk section names the three most likely failure points and the mitigation for each.

The 68 prompts

Every prompt below is written to produce a **finished deliverable**, not a list of ideas. Each one names the artifact, its sections, its row counts, and its format. Each one includes the placeholder list, the handling rules, and the enforcement of `{{brand_voice_guide}}` where client-facing words are produced.

How to read a prompt

- The **Role** line sets the standard. Change it if your client needs a different register.
- The **Inputs** block is what you paste in. Every placeholder must be replaced.
- The **Output format** block is the contract. Do not remove it — it is the reason the output arrives in a usable shape.
- The **Rules** block prevents the three most common failures: generic copy, fabricated data, and wrong structure.
- The **How to use** line after each prompt tells you what to check before the artifact goes anywhere near a client.

Prompts are numbered continuously from 01 to 68 so you can reference them in your Notion database, your team's Slack, or your client proposal ("the strategy brief is Prompt 03, run against the discovery transcript").

Client Strategy & Discovery

01 — Discovery Call Synthesis Into a Client Brief

```
ROLE:
You are a senior strategist at {{your_agency}}. Your job is to turn a raw discovery call into a
written brief a delivery team can build from, without a follow-up call.

INPUTS:
- Client: {{client_name}}
- Industry: {{industry}}
- Discovery transcript or notes: {{discovery_notes}}
- Brand voice guide: {{brand_voice_guide}}
- Known goals: {{goals}}

OUTPUT FORMAT — produce exactly these five sections, in this order:

1. CLIENT SNAPSHOT (4 bullets, max 60 words total): what they sell, who they sell it to, their
current marketing setup, the relationship as it stands today.
2. STATED PROBLEM (3 bullets): the problems the client named, in their own words, quoted where
possible.
3. UNDERLYING PROBLEM (1 paragraph, 80–120 words): the pattern beneath what they said — the real
constraint holding results back.
4. GOALS & CONSTRAINTS (a table with columns: Goal | Success metric | Deadline | Constraint).
Minimum 3 goal rows, maximum 6. If the client gave no metric, write [NEEDS INPUT: metric for
{{goal}}] in that cell — do not invent one.
5. OPEN QUESTIONS (a numbered list, 5–8 items): only questions that materially change what we
deliver. For each, add one line: "Why it matters: {{impact}}".

RULES:
- Separate what the client SAID from what you INFER. Label inferences as "Hypothesis".
- Never invent a metric, deadline, budget, or customer name. Use [NEEDS INPUT: ...] instead.
- Write in plain professional English. No consultancy vocabulary ("leverage", "synergy", "unlock",
"holistic").
- If {{brand_voice_guide}} is provided, write any client-facing summary paragraphs in that voice;
the internal sections stay neutral.
- Do not recommend tactics in this brief. This document defines the problem, not the solution.

HOW TO USE:
Run this immediately after the first discovery call, before you write anything else. The "Open
Questions" list becomes your follow-up email — send it the same day. If the goals table is full of
[NEEDS INPUT] markers, do not start strategy work until at least two are closed.
```

02 — Client Goals Into a Measurable KPI Framework

```
ROLE:
You are a performance marketing strategist translating vague business goals into a KPI framework a client will recognise and a team can report against.

INPUTS:
- Client: {{client_name}}
- Business goals: {{goals}}
- Current metrics: {{current_metrics}}
- Industry benchmarks: {{industry_benchmarks}}
- Paid, organic, and owned channels in scope: {{channel_mix}}

OUTPUT FORMAT — produce three tables and one short paragraph:

PARAGRAPH (80 words): the logic connecting the business goal to the measurable behaviour that signals progress. Name the chain explicitly: "Because the business needs {{goal}}, we need {{behaviour}}, which we measure by {{metric}}."

TABLE 1 — KPI FRAMEWORK
Columns: Business goal | Objective | Primary KPI | Secondary KPIs | Current baseline | 30-day target | 90-day target | Source of truth
Rows: one per business goal. Minimum 3, maximum 6.

TABLE 2 — FUNNEL MAPPING
Columns: Funnel stage | User question | Behaviour we need | Metric | Where it is measured | Typical benchmark
Rows: exactly 5 (Awareness, Consideration, Conversion, Retention, Advocacy).

TABLE 3 — MEASUREMENT PLAN
Columns: Metric | Tool | Who pulls it | Cadence | Where it appears in reporting | Owner
Rows: one per KPI in Table 1, minimum 3 rows.

RULES:
- Every KPI must have a baseline. If no baseline exists, write "Baseline: [NEEDS INPUT — 30 days of data required]".
- Benchmarks must be labelled as benchmarks, not promises.
- Do not include vanity metrics in Table 1 (no follower counts, no impressions alone). If a vanity metric is genuinely important to the client, put it in Table 2 only.
- Prefer one primary KPI per goal. Multiple primaries means no primary.

HOW TO USE:
Use this before the strategy one-pager (Template 7). If the client pushes back on a target, ask what they think the target should be — then re-run with their number and record the disagreement in the notes. A target the client agrees with is worth more than a target you believe in.
```

03 — Ideal Customer Profile From Real Sales Data

```
ROLE:
You are a research analyst building an ideal customer profile from evidence, not assumptions.

INPUTS:
- Client: {{client_name}}
- Product or service: {{offer}}
- Deal notes, CRM fields, or closed-won interview notes: {{customer_data}}
- Churn reasons or lost-deal notes: {{lost_deal_data}}
- Brand voice guide: {{brand_voice_guide}}

OUTPUT FORMAT — produce these sections:

1. ICP IN ONE SENTENCE (max 30 words): who buys, what situation triggers the purchase, and what outcome they want.
2. ICP TABLE — Columns: Attribute | Best-fit customer | Median customer | Poor-fit customer. Rows: exactly 8 attributes covering company size, industry, role seniority, trigger event, budget band, current alternative, decision process, and primary objection. Fill with specific values, not adjectives alone.
3. SIGNALS TABLE — Columns: Signal | Where we observe it | How strongly it predicts fit | Action. Minimum 8 rows. Rank the top 3 by "how strongly it predicts fit".
4. ANTI-ICP (3 bullets): who to disqualify, and the cost of taking them on.
5. MESSAGE IMPLICATIONS (5 bullets): what to say to the best-fit customer that would be wrong to say to everyone.

RULES:
- Base every attribute on evidence in the input. If the input is thin, state "Evidence gap: {{attribute}}" and write the provisional read with a question mark.
- Use revenue or deal-size data when provided; otherwise use frequency.
- Do not describe a demographic profile unless the data supports it. Firmographics and triggers beat demographics.
- The client-facing message implications must respect {{brand_voice_guide}} — write them in the voice the best-fit customer would respond to.

HOW TO USE:
This is the input to every persona sheet (Template 14) and to the targeting prompts in Paid Advertising. If your CRM data is thin, this prompt still works — it will surface the gaps, and those gaps are often the highest-value output. Bring them to the client as questions, not as findings.
```

04 — Discovery Questionnaire That Pre-Sells the Engagement

```
ROLE:
You are a strategist preparing a client discovery questionnaire that produces usable answers. A good questionnaire is also the first proof of competence.

INPUTS:
- Client: {{client_name}}
- Engagement type: {{engagement_type}} (retainer, project, audit, campaign)
- What we need to know to deliver: {{information_requirements}}
- Known context from the referral or pitch: {{known_context}}
- Brand voice guide: {{brand_voice_guide}}

OUTPUT FORMAT — produce:

1. INTRODUCTION (60–80 words): sent to the client, written in {{brand_voice_guide}}. States why we're asking, what it unlocks, how long it takes, and that their answers shape the deliverable — never "to better understand your business."
2. QUESTION TABLE — Columns: # | Section | Question | Why we ask | Answer type | Required? | What we'll do with it. 24 questions across 6 sections: Business & goals, Audience & offer, Current channels & performance, Brand & creative, Measurement & tooling, Process & constraints. Answer type is one of: short text, long text, number, date, multiple choice, file reference, "they may not know".
3. "I DON'T KNOW YET" SECTION — 8 items: the things a client honestly may not know yet, with a sentence on how we'll answer them for them.
4. REQUESTED MATERIALS CHECKLIST — 12 items: files, access, and exports to request, with the format needed (CSV, PDF, PNG, screen recording).
5. DEADLINE & FOLLOW-UP NOTE — 2 sentences: when we need it and what happens next.

RULES:
- No question can be answered with "yes/no" unless a yes/no answer is genuinely actionable.
- No question asks about the client's preferences rather than their behaviour. "What tone do you like?" is banned; "Send me three emails you're proud of" is right.
- Ask for numbers with a date range attached: "Revenue for the last 6 months" not "revenue".
- Maximum 8 minutes of client effort. If the questionnaire exceeds that, cut questions.
- The "I don't know yet" section is written kindly — clients are busy, not negligent.

HOW TO USE:
Send this 48 hours after the kick-off call, not at it. When it comes back, run Prompt 01 against the answers and the received files together — you now have a brief built from their words, which is a far better foundation than pitch notes.
```

05 — Strategy Workshop Agenda and Facilitator Notes

```
ROLE:
You are a strategist running a 90-minute strategy workshop. You need an agenda that produces
decisions, not a discussion that produces notes.

INPUTS:
- Client: {{client_name}}
- Attendees and roles: {{attendees}}
- Pre-work outputs to build on: {{discovery_brief}}
- Time available: {{workshop_duration}}
- Decisions that must be made in the session: {{decisions_required}}

OUTPUT FORMAT — produce:

1. WORKSHOP OBJECTIVE (1 sentence): what is true at the end of this session that is not true now.
2. AGENDA TABLE — Columns: Time | Segment | Objective | Facilitator prompt (exact words to say) |
Format | Output captured. Rows sum exactly to {{workshop_duration}}. Include one 10-minute break.
3. EXERCISE CARDS — 3 cards, each with: name, purpose, step-by-step instructions, the exact
question written on the whiteboard, time limit, and what the facilitator does if the group goes
quiet.
4. DECISION LOG — pre-filled with a row per decision in {{decisions_required}}: Decision | Options
on the table | Who decides | Deadline if not decided today.
5. PARKING LOT — a list structure for off-topic items, with a line reading "We'll come back to this
— logged here."
6. CLOSING SCRIPT (80 words) and a FOLLOW-UP EMAIL (120 words) that goes out within 24 hours: what
was decided, what remains open, what we need from the client and by when.

RULES:
- Every agenda row states an output captured. If a segment produces nothing, cut it.
- Facilitator prompts are actual sentences to say, including the question that unblocks a quiet
room.
- No exercise exceeds 20 minutes.
- Decision log rows must have a named decision-maker, not "the team".
- The follow-up email contains no new analysis — only what was said in the room.

HOW TO USE:
Print the agenda and the decision log; keep the follow-up email open in a second window. Send it
within 24 hours — the strategic value of a workshop decays fast, and the client will read that
email as proof the session was worth their time.
```

06 — Positioning Statement and Messaging Hierarchy

```
ROLE:
You are a brand strategist producing a positioning statement and the messaging hierarchy that sits under it.

INPUTS:
- Client: {{client_name}}
- Product or service: {{offer}}
- Audience research: {{audience}}
- Competitors: {{competitors}}
- Brand voice guide: {{brand_voice_guide}}
- Claims we can make (legal-approved): {{approved_claims}}

OUTPUT FORMAT:

1. POSITIONING STATEMENT (max 45 words): "For {{audience segment}} who {{need}}, {{client_name}} is the {{category}} that {{benefit}} because {{reason to believe}}." Fill every blank. If a blank cannot be filled from the inputs, write [NEEDS INPUT: ...] inside it — do not fill it with a safe-sounding phrase.
2. CATEGORY & FRAME-OF-REFERENCE (1 paragraph, 60 words): the category the client should be understood in, and the existing frame the strategy borrows from.
3. MESSAGING HIERARCHY — Columns: Tier | Message | Audience relevance | Proof point | Where it's used | Words to avoid. Rows: exactly 3 tiers — Tier 1 primary message (1 row), Tier 2 supporting messages (3 rows), Tier 3 proof points (4 rows).
4. DIFFERENTIATION TABLE — Columns: Claim | Competitor position | Our position | Substantiation. Minimum 4 rows. At least 2 rows must be honest about where a competitor is stronger.
5. ELEVATOR PITCH (3 versions): 15 words, 30 words, 60 words.
6. WHAT WE DO NOT SAY (4 bullets): claims to avoid, with the reason.

RULES:
- Every tier-1 and tier-2 message must trace to an approved claim. If {{approved_claims}} is missing, write [NEEDS INPUT: legal-approved claims] at the top of the hierarchy table and mark all rows as provisional.
- A message that could be used by any competitor is not a differentiator — flag it.
- Nothing in the hierarchy may contradict {{brand_voice_guide}}.
- Write the elevator pitches in the brand voice; they are spoken out loud.

HOW TO USE:
This output feeds landing page copy (Prompt 44), ad copy (Prompt 23), the strategy one-pager (Template 7), and proposal covers. Run it once, store it, and treat it as the source of truth — every later piece of client-facing copy should trace back to the hierarchy in this document.
```

07 — Competitive Landscape Scan for a Pitch

```
ROLE:
You are a strategist building a competitive scan that supports a pitch. It must be defensible, current, and useful on the slide — not a Wikipedia summary of the category.

INPUTS:
- Client: {{client_name}}
- Category and market: {{industry}}, {{market}}
- Named competitors: {{competitors}}
- The client's offer: {{offer}}
- Public information you have been given: {{competitor_research}}
- Brand voice guide: {{brand_voice_guide}}

OUTPUT FORMAT — produce:

1. MARKET STRUCTURE (150 words): who the buyer actually compares against, including indirect competitors and the alternative of doing nothing.
2. COMPETITOR MATRIX — Columns: Competitor | Positioning line | Target segment | Offer & price | Primary channel | Signature asset | Visible weakness. Rows: exactly 5 competitors.
3. THE WHITE SPACE (1 paragraph, 80 words): the territory no competitor occupies, and the reason it's worth occupying.
4. THE CLIENT'S FIT (1 paragraph): where {{client_name}} genuinely fits that white space, and where it doesn't. The second half is not optional.
5. PITCH ANGLE — 3 options: Columns: Angle | The one-line argument | Proof we'd need | Risk. Include one angle that is contrarian rather than safe.
6. SOURCES TABLE — Columns: Claim | Source | Date checked | Confidence. Every factual claim in sections 1–4 appears here.

RULES:
- Never state a competitor's price, revenue, or strategy without a source and a date. If you don't have it, write [NEEDS INPUT: verify {{competitor}} pricing].
- Do not describe competitors as "bad". Describe them as differently positioned. If a competitor is genuinely stronger on something, say so — clients respect it and it makes the rest of the analysis credible.
- The white space must be reachable by this client. A white space in a segment they can't serve is decoration.
- Pitch angles follow {{brand_voice_guide}} and are written to be said out loud in a room.

HOW TO USE:
Run this 48 hours before a pitch, then verify every claim in the sources table yourself — models get pricing and positioning details wrong, and a client will check. Replace anything you can't verify with a public source or cut it.
```

08 — Business Case and ROI Model for a Retainer

```
ROLE:
You are a commercial strategist building the business case for a retainer engagement, using the client's own numbers where you have them.

INPUTS:
- Client: {{client_name}}
- Proposed scope and fee: {{scope}}, {{fee}}
- Current performance baseline: {{current_metrics}}
- Client's cost or lead value: {{unit_economics}}
- Team hours included per month: {{hours_included}}
- Brand voice guide: {{brand_voice_guide}}

OUTPUT FORMAT — produce:

1. THE ARGUMENT (120 words): why this engagement is the cheapest available path to the client's goal. Written in plain language for a business owner, not a media buyer.
2. ROI MODEL TABLE — Columns: Line | Value | Source / assumption. Rows: fee, hours included, blended hourly cost, current monthly output (leads/sales/content), value per unit, projected monthly output, projected monthly value, net monthly value, ROI multiple. Every number carries a source or is labelled "ASSUMPTION — client to confirm".
3. SCENARIO TABLE — Columns: Scenario | Output assumption | Monthly value | Net vs fee | Break-even requirement. Exactly 3 rows: conservative, expected, stretch.
4. RISK & DOWNSIDE (5 bullets): what makes this not work, and the trigger that would prompt a scope conversation.
5. WHAT WE'LL MEASURE (table, 4 rows): Metric | Current | Target | How we'll know we're on track.
6. THE ASK (60 words): what we need from the client to make the model true — access, decision speed, content approval turnaround.

RULES:
- Every number is either sourced from the inputs or labelled ASSUMPTION. No blended averages that hide their inputs.
- If the ROI is negative on the conservative scenario, say so plainly and explain what changes it. A model that only works on the best case is not a business case.
- Never use "guaranteed", "risk-free", or "triple your revenue".
- The argument paragraph must survive being read aloud by the client's CFO.

HOW TO USE:
Put this in the proposal as an appendix, not as the opening slide — clients trust the appendix more than the pitch. If they challenge a number, fix the assumption and re-run; never defend a soft number.
```

Content Production

09 — Long-Form Article From a Brief

```
ROLE:
You are a senior content writer producing publication-ready copy for {{client_name}}. The output must be publishable without a rewrite.

INPUTS:
- Topic and target keyword: {{topic}}, {{target_keyword}}
- Search intent and audience: {{search_intent}}, {{audience}}
- Client offer and key proof points: {{offer}}, {{proof_points}}
- Brand voice guide: {{brand_voice_guide}}
- Approved claims and claims to avoid: {{approved_claims}}, {{claims_to_avoid}}
- Internal links available: {{internal_links}}
- Target word count: {{word_count}}

OUTPUT FORMAT — produce these sections in this order:

1. TITLE (2 options, max 60 chars each, then the reason each was chosen in 10 words)
2. META DESCRIPTION (1, 120–155 chars, includes the target keyword once, includes a reason to click)
3. SLUG (lowercase, hyphenated, max 5 words)
4. OPENING (2 paragraphs, 120 words): the reader's problem stated in their words, then the promise of the article. No throat-clearing.
5. H2/H3 OUTLINE (as a markdown list): every H2 and H3 with a one-line note on what that section contains. Minimum 4 H2s.
6. FULL DRAFT: the complete article under those headings, with all paragraphs, lists, and the final CTA. Target {{word_count}} ±10%. Do not write "..." or "[continue]" — write the whole thing.
7. CTA BLOCK (final 60 words): the ask, plus the reason to act now, written in the client's voice.
8. INTERNAL LINK PLAN (table): Anchor text | Target URL | Placement in the article
9. EDITOR'S NOTES (4 bullets): anything the client must verify, plus 2 places where a real number, screenshot, or example would strengthen the piece.

RULES:
- Match {{brand_voice_guide}} strictly: sentence rhythm, vocabulary, formality level, and banned words.
- No invented statistics, studies, customer names, or quotes. If a number is needed, write [CLIENT STAT: ...] in the copy exactly where it goes.
- Write for a human reader, not for a crawler. No keyword stuffing; use the target phrase naturally where a person would.
- Never use "in today's digital landscape", "it's important to note", "delve into", "unlock", "leverage", "game-changer", "revolutionise".
- If the outline would require more than {{word_count}} words, compress the outline rather than truncating sentences.

HOW TO USE:
Generate, then run Prompt 57 (Editorial QA Pass) over the draft before it goes to the client. Treat the "Editor's notes" section as your list of client questions — it converts a vague "can you review this?" into a specific list of three things to check.
```

10 — Blog Post Series From One Pillar Topic

```
ROLE:
You are a content strategist turning one pillar topic into a coherent series that builds authority and internal links.

INPUTS:
- Pillar topic: {{pillar_topic}}
- Target keyword and supporting keywords: {{target_keyword}}, {{supporting_keywords}}
- Audience and their stage: {{audience}}
- Brand voice guide: {{brand_voice_guide}}
- Proof points and offers to reference: {{proof_points}}, {{offer}}
- Pillar page URL (for internal links): {{pillar_url}}

OUTPUT FORMAT:

1. SERIES OVERVIEW (100 words): the argument the series makes in total, and what the reader can do after finishing it.
2. SERIES PLAN TABLE — Columns: # | Working title | Target keyword | Intent | Angle | Length | Publish week | Links to | Funnel stage. Exactly 6 posts.
3. FULL OUTLINES — for each of the 6 posts: H1, meta description, 4–6 H2s with a sentence of content under each, the proof point used, the CTA, and the posts in the series it links to.
4. LINK GRAPH (table): From post | To post | Anchor text | Why this link
5. HOOK BANK — 18 opening lines across the 6 posts, grouped by post, each written as a finished sentence. Two per post plus three alternates.
6. REPUBLISHING PLAN (table): Post | Reuse as | New format | Timing | Channel

RULES:
- Each post must be independently useful and must also be incomplete alone — the series has a reason to exist.
- The angles must not overlap. If two posts argue the same point, merge or re-angle them.
- Only link to a post that will exist; mark planned links as [planned].
- Every hook must be specific to its post. Generic hooks are banned.
- Apply {{brand_voice_guide}} to every title, hook, and CTA.

HOW TO USE:
Use the series plan as a client-facing roadmap in a content meeting — clients approve direction faster when they see six concrete titles rather than a topic list. Lock the plan before writing a single word of draft.
```

11 — Case Study With a Real Result

```
ROLE:
You are a writer turning project material into a case study that a client can publish with their name on it.

INPUTS:
- Client: {{client_name}}
- Subject company: {{subject_company}} (or "our client" if permission is pending)
- Challenge and starting position: {{challenge}}
- What we did: {{approach}}
- Results with numbers: {{results}}
- Testimonials available: {{testimonials}}
- Brand voice guide: {{brand_voice_guide}}
- Permission status for naming and logos: {{permissions}}

OUTPUT FORMAT — produce these sections:

1. HEADLINE (2 options, under 60 chars, each with the result in it)
2. META DESCRIPTION (1, 140–155 chars)
3. AT A GLANCE (table): Client | Sector | Challenge | Solution | Result | Timeframe. One row.
4. THE CHALLENGE (120 words): specific, quantified, no drama. What was broken or missing, with the number that proved it.
5. THE APPROACH (180 words, 3 sub-headings): what we did, in the order we did it. Concrete enough that another team could repeat it.
6. THE RESULTS (table): Metric | Before | After | Change | Timeframe. Minimum 3 metric rows, maximum 6. Also include 1–2 sentences of interpretation per metric.
7. QUOTE FROM THE CLIENT (formatted, with name, role, company — or [CLIENT QUOTE: ...] if not yet approved)
8. KEY TAKEAWAYS (3 bullets): what a reader in a similar position should do, phrased as advice.
9. CTA (40 words)
10. PUBLISHING NOTES (4 bullets): permission dependencies, assets needed (logo, screenshot, headshot), and the recommended placement on the site.

RULES:
- Results must come from {{results}}. Never round up, never extrapolate, never combine a 6-month figure with a 3-month timeframe.
- If a metric has no before-value, write "Before: not tracked" — don't manufacture a baseline.
- Write in {{brand_voice_guide}}, and prefer plain claims to superlatives.
- If {{permissions}} does not include naming rights, write the whole draft in the "our client" form and flag it in Publishing Notes.

HOW TO USE:
Send the draft with the results table highlighted — clients approve numbers fastest and slowest, so isolate them. Keep the "Before: not tracked" rows visible; they explain why some metrics can't be compared and they make the rest of the table believable.
```

12 — Product Description Set for an Ecommerce Page

```
ROLE:
You are an ecommerce copywriter producing the full description set for one product page.

INPUTS:
- Product: {{product_name}}
- Specs and materials: {{product_specs}}
- Who it's for and the problem it solves: {{audience}}, {{problem}}
- Price and what's included: {{price}}, {{includes}}
- Brand voice guide: {{brand_voice_guide}}
- Approved claims: {{approved_claims}}

OUTPUT FORMAT — produce:

1. PAGE TITLE (≤ 60 chars) and META DESCRIPTION (≤ 155 chars)
2. HERO BLOCK: H1 (≤ 10 words), subhead (20–30 words), 3 bullet value props (≤ 8 words each)
3. SHORT DESCRIPTION (2 sentences, ≤ 40 words) for product cards and search results
4. LONG DESCRIPTION — 5 sub-sections with headings: what it is, why it works, who it's for, what's
in the box, how to use it. 200–300 words total, written as final copy.
5. SPECIFICATION TABLE — Columns: Specification | Detail. 8–12 rows, including dimensions, weight,
materials, care, warranty, what's included.
6. BENEFITS TABLE — Columns: Feature | What it means for the customer | Use case. One row per
feature from {{product_specs}}.
7. USE-CASE ROW (table): Situation | How they use it | What they get. 4 rows.
8. FAQ (6 questions): Question in the customer's voice | Answer (40–60 words). Cover: fit, care,
compatibility, returns, gifting, comparison to alternatives.
9. CTA PAIR: Primary button text | Secondary link text | One supporting line
10. IMAGE REQUIREMENTS (table): Image # | What it shows | Format | Aspect | Priority — 6 rows for
the product page gallery.

RULES:
- Every claim must appear in {{approved_claims}} or be a factual spec. No "life-changing", no
"you'll never", no health or performance claims not documented.
- Benefits must be written for the customer's situation, not the product's features. "Holds 12
drinks" → "Fits a full day's water without refilling".
- Match {{brand_voice_guide}} precisely — product copy is where brand voice is most visible and
most often breaks.
- Never compare to a named competitor. Comparative claims without substantiation are a legal risk.

HOW TO USE:
Pair this with Prompt 55 (Image Prompt Set for Product Photography) and you can ship a product page
from a single spreadsheet of specs. Run Prompt 58 (Legal & Claims Check) before publishing anything
with a performance or health claim.
```

13 — Landing Page Headline and Hero Copy Set

```
ROLE:
You are a conversion copywriter producing the hero block and the headline set for a landing page, then the section headlines that follow.

INPUTS:
- Offer: {{offer}}
- Audience and their current belief: {{audience}}, {{current_belief}}
- Single-minded proposition from strategy: {{smp}}
- Proof available: {{proof_points}}
- Traffic source the page receives: {{traffic_source}}
- Brand voice guide: {{brand_voice_guide}}
- Claims to avoid: {{claims_to_avoid}}

OUTPUT FORMAT — produce:

1. ANGLES TESTED (table): Angle | Core promise | Proof anchor | Risk. Exactly 4 angles — outcome, pain relief, objection removal, and contrarian.
2. HERO BLOCK — for the top 2 angles only. For each: H1 (≤ 12 words), subhead (20–30 words), primary CTA (2–4 words), secondary CTA (2–4 words), trust line (≤ 15 words), and a one-line note on why this angle suits {{traffic_source}} traffic.
3. SECTION HEADLINE SET (table): Section number | Section purpose | Headline | Subhead | What proves it. Exactly 7 sections.
4. CTA MATRIX (table): Placement | Button text | Supporting microcopy | Expected action. 6 placements, with consistent button text across all rows.
5. VARIATION SET: 3 alternate H1s for the recommended angle, each testing a different mechanism.
6. REJECTION NOTES (4 bullets): the angle you are NOT recommending, and why — with what it would take for it to become the right choice.

RULES:
- The H1 states the outcome, not the mechanism. If a reader can't tell what they get, the headline is wrong.
- Match {{brand_voice_guide}} exactly — headline register matters more than body copy, and clients notice headlines first.
- Every headline must be paired with a proof anchor. A headline with nothing behind it is a promise the page can't keep.
- Do not use questions as headlines unless the question is exactly what the audience asks out loud.
- No fake urgency, no countdown pressure unless there is a real deadline.

HOW TO USE:
Send the four angles to the client and let them pick before you write a word of body copy — it is the cheapest decision in the whole project and it prevents a full rewrite. Once they choose, the section headline set maps directly onto Template 11.
```

14 — Video Script With Shot List

```
ROLE:
You are a video producer writing a production-ready script for {{client_name}}, with a shot list a small crew can shoot in one session.

INPUTS:
- Video type and platform: {{video_type}}, {{platform}}
- Length: {{target_length}}
- Objective: {{objective}}
- Audience and what they should do afterwards: {{audience}}, {{desired_action}}
- Message hierarchy and proof: {{message_hierarchy}}, {{proof_points}}
- Brand voice guide: {{brand_voice_guide}}
- Who is on camera: {{on_camera_talent}}

OUTPUT FORMAT — produce:

1. VIDEO SPEC (table): Platform | Length | Aspect ratio | Resolution | Captions required | CTA
2. HOOK OPTIONS (3 finished lines, each ≤ 12 words, with a note on the visual that plays under each)
3. SCRIPT TABLE — Columns: Timecode | Visual / shot | On-camera words (final) | On-screen text | Audio / music | Notes. Every row timed; the timecodes run continuously from 00:00 to the full length. Minimum 10 rows.
4. B-ROLL LIST (table): Shot | Description | Where it comes from | Priority. 8 rows.
5. CAPTIONS FILE — the full spoken script as a plain caption block, broken into lines of ≤ 42 characters, with speaker labels and punctuation for the burned-in version.
6. CTA AND END CARD (table): Element | Final content | Duration
7. PRODUCTION CHECKLIST (table): Item | Owner | Needed by | Status — 10 rows, including talent, location, props, wardrobe, releases, and platform access.
8. ALT TEXT (one line per shot): a description of the visual for accessibility.

RULES:
- The first 3 seconds must contain the hook. Write the words that go there explicitly.
- Spoken lines must be sayable at normal pace in the timecode given. If a line is too long, cut words, not time.
- Do not write "insert B-roll here" as a script line — write the shot.
- Apply {{brand_voice_guide}} to every spoken and on-screen word. Read it aloud before sending.
- If {{proof_points}} is empty, use a demonstration beat instead of a claim beat. Do not invent a result.

HOW TO USE:
The captions file is the single most reusable output here — send it to the client for approval even when they approve the script visually, because caption text is what most people actually read. Add shot list row numbers to the edit brief.
```

15 — Newsletter From a Single Idea

```
ROLE:
You are a writer producing a complete, sendable email newsletter from one idea — not a roundup of links.

INPUTS:
- The one idea: {{newsletter_idea}}
- Audience and what they care about: {{audience}}
- Supporting points and proof: {{supporting_points}}
- Client offer or CTA: {{offer}}, {{cta}}
- Brand voice guide: {{brand_voice_guide}}
- Previous subject lines that worked and failed: {{subject_line_history}}
- Sending metadata: {{send_date}}, {{segment}}

OUTPUT FORMAT — produce:

1. SUBJECT LINE OPTIONS (5): Text | Character count | Predicted angle | Why it works
2. PREHEADER (1, 40–90 chars, complements rather than repeats the subject)
3. SEND-TIMING RECOMMENDATION: 1 line with day, time, and timezone, plus the reasoning in one sentence
4. EMAIL BODY — 6 blocks, written as final copy: Hook (≤ 40 words), Setup (60 words), Idea explained (150 words), Proof or example (80 words), Practical takeaway (80 words, 3–5 bullets), CTA (40 words)
5. BUTTON TEXT + LINK (table): Placement | Button text | Link destination | Tracking note
6. PLAIN-TEXT VERSION (the full email with no formatting)
7. SEGMENT NOTES (3 bullets): who should get this and who should be excluded
8. QA CHECKLIST (table, 8 rows): Check | Pass condition | Result
9. FOLLOW-UP EMAIL (optional, 90 words): a 3-day follow-up for non-openers with a different subject and a different angle.

RULES:
- One idea per newsletter. If the input contains three, select one and note the other two as "next issue" in a single line.
- No link-heavy roundups. If a link is included, it supports the idea; it is not the idea.
- Match {{brand_voice_guide}} — email is the most personal channel and the easiest place to break voice.
- Every claim about the client's results must come from {{supporting_points}}. If you need a number, write [CLIENT STAT: ...].

HOW TO USE:
Pick the subject line from the five using your own open-rate history, not the model's prediction. Send the follow-up version to non-openers only — the second email has to be materially different, or it reads as nagging.
```

16 — Content Repurposing Matrix From One Long Asset

```
ROLE:
You are a content strategist turning one substantial asset into a full week of channel-native content, with each output written as final copy.

INPUTS:
- Source asset: {{source_asset}} (article, video, webinar, or podcast transcript)
- Channels in scope: {{channel_mix}}
- Audience: {{audience}}
- Brand voice guide: {{brand_voice_guide}}
- Constraints (offer links, approvals, embargoes): {{constraints}}

OUTPUT FORMAT — produce:

1. SOURCE DIGEST (100 words): the core argument, the three strongest moments, and the proof points available.
2. REPURPOSING MATRIX — Columns: # | Output | Channel | Format | Hook (finished first line) | Core message | Source moment used | CTA | Publish day | Effort. Exactly 10 rows — 7 content posts, 2 emails, 1 ad.
3. FULL COPY for each of the 10 rows, written as final copy in the correct platform format. Include character counts for social posts.
4. FORMAT NOTES (table): Output | Platform-specific requirements | Visual spec | Caption/alt text. One row per output.
5. SEQUENCING PLAN (table): Day | Output | Purpose in the sequence
6. WHAT TO SKIP (3 bullets): moments in the source that do not work as standalone content, with the reason.

RULES:
- Each output must be native to its platform. Do not paste the same paragraph into five channels with different hashtags.
- A post must be readable without clicking through. If it only makes sense with the source, it's a teaser — write a teaser only if that's the goal, and label it.
- Apply {{brand_voice_guide}} to every output, and adjust formality by channel without breaking the voice.
- No output may reuse the same hook as another.

HOW TO USE:
Run this the day a core asset is signed off, and schedule the whole matrix at once. The "what to skip" list saves an hour of arguing about whether a moment is strong enough — it is already decided.
```

Social Media

17 — 30-Day Social Content Calendar

```
ROLE:
You are a social media lead producing a client-ready 30-day content calendar for {{client_name}}.

INPUTS:
- Brand voice guide: {{brand_voice_guide}}
- Content pillars: {{content_pillars}}
- Channels: {{channel_mix}}
- Goals for the month: {{goals}}
- Offer and primary CTA: {{offer}}, {{primary_cta}}
- Seasonal or launch dates: {{key_dates}}
- Approval constraints: {{constraints}}

OUTPUT FORMAT — produce exactly:

HEADER BLOCK (a table): Client | Period (exact start and end dates) | Channels | Pillars | Month goal | Default CTA | Approval path

CALENDAR TABLE — Columns: Date | Day | Channel | Pillar | Format | Post type | Hook / first line | Full caption or body copy | CTA | Asset + owner.
Exactly 30 rows, one per calendar day, in date order starting from the first date in {{period}}. Weekend days are included; if a channel skips weekends, use "— (rest day)" in Channel and a reason in Post type.

RULES:
- Every one of the 30 rows has content in every cell. No blank cells, no "TBC", no "[placeholder]" left in the copy. If something genuinely can't be decided, write [NEEDS INPUT: ...] in that single cell and nowhere else.
- Hooks are finished sentences a person would say, not topic labels. "Benefits of CRM" is a label; "Your CRM is full of leads you'll never call back" is a hook.
- Captions are final copy, ready to paste. Do not outline.
- No two rows share the same hook or the same post type on the same channel within any 7-day window.
- Pillar distribution across the 30 days must match {{content_pillars}} weights (if not given, distribute roughly 40% education, 30% proof, 20% personality, 10% offer and say so in one line).
- Every caption follows {{brand_voice_guide}}. Do not use words the guide bans, even as filler.
- The final row is always a recap or CTA post.

HOW TO USE:
Send this as the monthly planning artefact with the captions already written — clients approve calendars faster when they read the copy rather than imagine it. Check the last row is a CTA, then paste the table straight into Sheets and convert to a Notion database.
```

18 — Weekly Content Calendar for Fast Approval

```
ROLE:
You are a social lead producing a single-week calendar designed to be approved on a phone in under 10 minutes.

INPUTS:
- Week starting: {{week_starting}}
- Brand voice guide: {{brand_voice_guide}}
- Channels: {{channel_mix}}
- Theme or campaign focus: {{week_theme}}
- The week's goal: {{week_goal}}
- Asset availability: {{assets_available}}
- Any hard dates: {{key_dates}}

OUTPUT FORMAT — produce:

WEEK SUMMARY (table): Week of | Theme | Goal | Reach-or-depth focus | Reused-from asset | Total assets due

CALENDAR TABLE — Columns: Day | Channel | Post type | Pillar | Hook / first line | Final caption or body copy | CTA | Asset + due date.
Exactly 7 rows, Monday to Sunday, in that order.

PRODUCTION NOTES (table): Asset | Owner | Format | Due day | Dependency — one row per asset in the calendar, 3–8 rows.

RULES:
- Exactly 7 rows. A rest day uses an em dash in every cell and states the reason in Post type ("Rest — batch Monday's assets"), never a silent blank.
- Captions are written final. No bullet outlines of what the post should say.
- Keep every cell short enough to read on a phone: hooks under 12 words, captions under 150 words.
- If a launch or live event falls in the week, mark that row's Day with a 📌.
- Apply {{brand_voice_guide}} to every hook and caption.
- Total assets in the Production Notes table must equal the number of non-rest days plus any extras listed.

HOW TO USE:
Send this on Friday afternoon for the following week, so the client approves it before Monday. If they ask for changes, edit the caption cell directly — the structure rarely needs to change, and rebuilding the table is where hours disappear.
```

19 — One Post, Five Platforms, Written Five Ways

```
ROLE:
You are a social media writer adapting a single idea into platform-native posts. Each version must feel written for its platform, not reformatted from a master copy.

INPUTS:
- The idea: {{content_idea}}
- Platforms: {{platforms}}
- Audience: {{audience}}
- Brand voice guide: {{brand_voice_guide}}
- Links or CTA destination: {{cta_destination}}
- Proof or detail to include: {{proof_point}}
- What worked before on each platform: {{platform_performance_notes}}

OUTPUT FORMAT — for each platform in {{platforms}}, produce:

PLATFORM BLOCK:
- Post type and length: (e.g. "LinkedIn text post, 130–180 words")
- Post copy (final, complete, ready to paste)
- Hook / first line (called out separately)
- Hashtags (only where the platform uses them; 3–5 max on LinkedIn, 0–3 elsewhere)
- Visual or video direction: one line
- Alt text: one line
- CTA placement: where and how
- Character count of the post copy

RULES:
- Never paste the same paragraph into two platforms. Different platform, different opening line, different rhythm, different call to action.
- LinkedIn: no emoji stacks, no "🚀 Excited to share", first line must work before "see more". TikTok: hook must land in the first 2 seconds, so the first line is a spoken line. Instagram: caption opens with the hook line, then a blank line, then the body. X: under 280 characters, no thread unless asked.
- Use {{proof_point}} as a concrete detail in every version.
- Apply {{brand_voice_guide}}, but adjust formality to the platform without breaking the voice.
- No hashtag stuffing, no engagement bait ("comment YES below"), no "link in bio" unless the destination is bio.

HOW TO USE:
Use this when a client wants "one idea, everywhere" without five times the cost. Check the LinkedIn first line against a phone preview — if it doesn't earn the tap, rewrite the first line before sending.
```

20 — Instagram Carousel: Slides and Caption

```
ROLE:
You are a designer-copywriter producing a complete Instagram carousel: the copy for every slide, the caption, and the visual instruction for each slide.

INPUTS:
- Topic: {{topic}}
- Goal: {{goal}}
- Number of slides: {{slide_count}}
- Audience: {{audience}}
- Brand voice guide: {{brand_voice_guide}}
- Proof points: {{proof_points}}
- Brand colours and typography: {{brand_assets}}
- CTA destination: {{cta_destination}}

OUTPUT FORMAT — produce:

1. CAROUSEL STRUCTURE TABLE — Columns: Slide # | Role | Slide headline (final) | Slide body (final, ≤ 35 words) | Visual instruction | Design notes. Exactly {{slide_count}} rows. Slide 1 is the cover; the last slide is the CTA.
2. COVER SLIDE — 3 headline options, each ≤ 9 words, with the reason each works.
3. FULL SLIDE COPY — every slide's headline and body written out as final copy, ready for layout.
4. CAPTION — 120–180 words, opening with the cover headline, then the value, then the CTA. Include a first-comment block of 3 extra lines.
5. CTA OPTIONS (3): Button text | Destination | Why this wording
6. ALT TEXT for each slide (one line each).
7. DESIGN HANDOFF NOTES (table): Element | Specification — typeface, sizes, colour hexes, margins, safe zones, 4:5 aspect, 1080×1350 px.
8. ENGAGEMENT PROMPT (1 line): a question that invites a real reply.

RULES:
- One idea per slide, except the cover and CTA. If a slide needs two ideas, split it and add a slide.
- Slide headlines under 9 words; body under 35 words. If a slide needs more, it's a caption paragraph, not a slide.
- Slide 1 must work standalone — most people see only the cover before swiping.
- No "swipe to see more" clichés. No numbered-list abuse like "5 Ways to...".
- Apply {{brand_voice_guide}} to every word on every slide. This is where voice guidelines usually get ignored — check it.
- The CTA slide has exactly one action.

HOW TO USE:
Send the slide table to the designer and the caption separately to the client approver — designers care about layout, clients care about words, and combining them slows both. Add the alt text to the upload step; it's a spec, not a nicety.
```

21 — Short-Form Video Scripts for Reels, TikTok, Shorts

```
ROLE:
You are a video writer producing three ready-to-shoot short-form video scripts for {{client_name}}, each with a spoken hook, a full script, shot notes, and captions.

INPUTS:
- Topic: {{topic}}
- Goal: {{goal}}
- Target duration: {{duration}}
- Platform: {{platform}}
- Audience: {{audience}}
- Brand voice guide: {{brand_voice_guide}}
- Proof or demonstration available: {{proof_point}}
- Who appears on camera: {{talent}}
- B-roll or product access available: {{assets_available}}

OUTPUT FORMAT — produce three scripts, each with:

SCRIPT BLOCK:
- Concept name and the one-line idea
- Hook (first 2 seconds, ≤ 9 spoken words, written as it is spoken)
- Spoken script with timecodes (0–2s hook, 2–5s setup, 5–{{duration}}s substance, final 2s CTA). Full sentences, not notes.
- On-screen text overlays with timing
- Shot list (table): Time | Shot | What happens | Notes
- Caption (≤ 90 words, opens with the hook)
- Hashtags (3–5, platform-appropriate)
- Alt text and a 1-line accessibility description
- CTA button text and destination
- Variants: 2 alternate hooks for A/B testing

RULES:
- The hook must create a reason to keep watching: a specific number, a contradiction, a mistake callout, or a promise with a timeframe.
- Write for the ear. Short sentences, no subordinate clauses stacked, no jargon the viewer wouldn't use.
- Every script has a visible or verbal proof moment — a screen recording, a result, a demonstration, or a physical object.
- Apply {{brand_voice_guide}} to spoken lines; check that the guide's formality survives being spoken aloud.
- No video ends without a spoken or on-screen CTA.
- Do not use trending audio references you can't verify — specify the audio type ("upbeat, mid-tempo, no vocals") instead.

HOW TO USE:
Shoot all three in one session and cut them later — that alone turns a 3-day production into a half-day one. Have the creator read each script aloud before recording; if they stumble, the line is too long.
```

22 — Community Management Reply Bank

```
ROLE:
You are a community manager writing ready-to-send replies to predictable audience comments, in the client's voice, with the line between helpful and over-promotional marked.

INPUTS:
- Brand voice guide: {{brand_voice_guide}}
- Platform: {{platform}}
- Common comment types: {{common_comments}}
- FAQ or approved answers: {{faq_approved_answers}}
- Offers and campaigns currently running: {{current_offers}}
- Escalation rules: {{escalation_rules}}

OUTPUT FORMAT — produce:

1. REPLY BANK TABLE — Columns: Comment type | Example comment | Reply (final, ready to paste) | Tone | Notes. Minimum 20 rows covering: genuine praise, price question, feature question, complaint, competitor mention, "how much?", "do you ship to X?", scam/spam, "can I get a discount?", job application, partnership request, request for free advice, negative review, wrong assumption about the product, and 6 more realistic cases.
2. ESCALATION TABLE — Columns: Situation | Response to send now | Internal action | Who is notified | Timeframe. 8 rows, including anything involving a legal threat, a medical/safety claim, a data request, or an influencer ask.
3. TONE LADDER TABLE — Columns: Situation | Temperature | Rule — 5 rows from warm to firm, describing when to move down the ladder.
4. PROACTIVE FAQ POSTS — 6 short posts (60–90 words each) answering the most-asked questions before they're asked, with a title for each.
5. NEVER-SAY LIST — 10 things not to say in replies, each with a one-line reason and a replacement phrase.

RULES:
- No reply may contain an unverified claim, a discount not in {{current_offers}}, or a promise not in {{faq_approved_answers}}.
- Never argue with a complaint. Acknowledge once, move to a DM, and escalate.
- Replies under 40 words for most cases; the reply must fit on a phone screen above the fold.
- Apply {{brand_voice_guide}}. Keep replies conversational even if the brand voice guide is formal elsewhere — and if it says so explicitly, follow the guide.
- Never use an emoji unless {{brand_voice_guide}} permits it and specifies which ones.

HOW TO USE:
Save this as a shared doc with a searchable table and let the community team copy-paste. After 30 days, log which reply types got the best response and update the bank — the fastest replies should be the ones you use most.
```

23 — Image and Graphic Prompt Set for Design and AI Generation

```
ROLE:
You are an art director producing a production-ready set of creative briefs for visual assets — one brief per asset, usable by a designer or pasted into an image generation tool.

INPUTS:
- Campaign or content theme: {{theme}}
- Number of assets needed: {{asset_count}}
- Brand colours, fonts, and visual rules: {{brand_assets}}
- Brand voice guide: {{brand_voice_guide}}
- Where assets will run: {{channels}}
- What the assets must communicate: {{message}}
- Product or subject: {{subject}}
- Photography available: {{available_photos}}

OUTPUT FORMAT — produce:

1. CREATIVE DIRECTION (120 words): the visual world, the feeling, the one visual idea that ties the set together, and the reference style described in words (not "modern and sleek").
2. ASSET BRIEFS — one block per asset, exactly {{asset_count}} blocks. Each block contains:
   - Asset ID (e.g. A01) and where it runs (channel, placement, dimensions)
   - The message this asset carries (one line)
   - Composition (what's in frame, where, what's cropped, camera angle)
   - Subject and action
   - Colour and light
   - Typography (if any) — exact text to appear
   - Do / don't list (3 each)
   - Alt text (one line)
   - AI generation prompt (a single paragraph, if applicable): subject, action, composition, lighting, lens, colour, style, aspect ratio, negative prompt
   - Design note (what a designer should watch out for)
3. CONSISTENCY RULES (table): Element | Rule | Applies to assets — 8 rows covering palette, lighting, model casting, typography, props, crop ratio, background, and treatment.
4. HANDOFF CHECKLIST (table): Check | Pass condition — 10 rows.

RULES:
- Describe the composition physically: what's in frame, left/right/top/bottom, what's cropped. "Left third, subject looking into the frame, negative space on the right for a 40-character headline" — that level of detail.
- Do not reference real people's likenesses or living artists by name.
- The AI generation prompt must include an aspect ratio and a negative prompt.
- Any text that appears in the image is written exactly and flagged for the designer to set as live type (never baked into a generated image).
- Apply {{brand_voice_guide}} to any written words in the design.
- Every asset has alt text.

HOW TO USE:
Run this with the creative brief (Prompt 48) in the same session so the words in the image and the words on the page agree. Designers will thank you for specifying the crop ratio — that's the reason images come back the wrong shape.
```

Paid Advertising

24 — Full Paid Media Account Build Plan

```
ROLE:
You are a performance media planner producing the account build plan for a new {{client_name}}
campaign — the document a media buyer executes from, without further interpretation.

INPUTS:
- Campaign objective: {{objective}}
- Offer: {{offer}}
- Target audience: {{audience}}
- Budget and split: {{budget}}, {{channel_mix}}
- Markets: {{market}}
- Tracking stack available: {{tracking_stack}}
- Landing page or offer URL: {{landing_url}}
- Brand voice guide: {{brand_voice_guide}}
- Constraints (brand safety, legal, seasonality): {{constraints}}

OUTPUT FORMAT — produce:

1. CAMPAIGN ARCHITECTURE TABLE — Columns: Campaign | Platform | Funnel stage | Objective | Audience
| Budget | Daily cap | Target CPA / ROAS | Bid strategy | Launch date. 6–8 rows.
2. AD SET / AD GROUP STRUCTURE — a nested table per campaign: Campaign | Ad set or group | Audience
definition (full targeting spec) | Exclusions | Placement | Budget | Bid | KPI. 12–18 rows total.
3. TRACKING SPECIFICATION — table: Event | Trigger | Parameter | Where it fires | Parameter value
format. Exactly 10 events covering view content, scroll 50%, form start, form submit, phone click,
add to cart, purchase, signup, download, and a qualified conversion.
4. UTM NAMING CONVENTION — table: Parameter | Format | Example. 5 rows, plus one full example URL
built with a real campaign name.
5. AUDIENCE & EXCLUSION RULES — 8 bullets covering exclusions, overlap priority, retargeting
windows, lookalike source and percentage, and suppression lists.
6. LAUNCH SEQUENCE — table: Date | Action | Owner | What we're watching | Rollback trigger. Minimum
6 rows across the first two weeks.
7. TESTING PLAN — table: Test | Hypothesis | Variable | Audience | Budget | Metric | Decision
threshold | End date. Minimum 6 rows.
8. BUDGET CONTROL RULES — 5 bullets: daily caps, the CPA at which we pause an ad set, the % drop
that triggers a check, the budget increase rule (e.g. +20% only after 3 days at target), and who
can approve spend changes.

RULES:
- Every targeting spec must be specific enough to build without asking a question: audience size
estimate, exclusions, placements, bid strategy.
- No target CPA without a basis. If there is no benchmark, write "Target CPA: [NEEDS INPUT — client
benchmark or 30-day test data]".
- The testing plan names a decision threshold in advance, in numbers, and a decision maker.
- Exclusions must include existing customers where that makes sense, plus employees and invalid
placements.
- Apply {{brand_voice_guide}} to any ad copy included in the plan and to the campaign naming
convention (use readable names, not codenames).

HOW TO USE:
This is the build doc you attach to the media buying ticket. Run it before opening a platform
account, not after — the structure in the table is what prevents the classic mess of 14 ad sets and
no learning data.
```

25 — Meta and Google Ad Copy Matrix

```
ROLE:
You are a performance copywriter producing a finished ad copy matrix — every cell contains final, platform-ready copy, not an angle.

INPUTS:
- Offer and promotion: {{offer}}, {{promotion}}
- Target audience: {{audience}}
- Unique selling points: {{usps}}
- Proof points: {{proof_points}}
- Brand voice guide: {{brand_voice_guide}}
- Landing page promise: {{landing_page_promise}}
- Claims to avoid: {{claims_to_avoid}}
- Competitor context: {{competitors}}

OUTPUT FORMAT — produce:

1. ANGLE SET TABLE — Columns: Angle | Core promise | Proof used | Audience tension | Risk of over-claiming. 6 rows, one per angle.
2. META MATRIX — Columns: Angle \ Format | Reel / video 9:16 | Carousel | Static / image. 6 angles × 3 formats = 18 cells. Each cell contains: Primary text (final, with character count), Headline (final, with character count), Description (final), CTA button, thumbnail direction, and the funnel stage it serves.
3. GOOGLE MATRIX — Columns: Ad type | Headlines (15, each with character count) | Descriptions (4, each with character count) | Path fields | Ad group. Rows: Responsive Search, Performance Max (headlines, long headlines, descriptions, 5 image prompts), YouTube (2 descriptions, 2 CTA variants), Display (2 short headlines, 1 long description).
4. ASSET PRODUCTION SHEET — Columns: Asset ID | Platform | Format | Spec | Copy | Design brief | Due | Status. One row per asset, 16–20 rows.
5. NEGATIVE KEYWORD & AUDIENCE LIST — 15 terms/phrases to exclude, each with the reason.
6. TESTING PLAN — table: Test | Hypothesis | Variable | Metric | Decision threshold | Decision date. 6 rows.

RULES:
- Character counts are shown next to every headline and description, and every one respects the platform limit.
- Every ad makes the same promise as the landing page it points to. If the ad over-promises, the page fails.
- The 18 Meta cells must be genuinely distinct. If two cells could run against the same audience at the same time without a learning problem, they are not different enough — rewrite one.
- Apply {{brand_voice_guide}} strictly. In paid copy, adjectives are the most common place a brand voice breaks.
- No "guaranteed", "risk-free", "best price", or "limited time" unless the deadline or claim is real and documented.
- No naming competitors in Meta copy. Comparative Google ads must be factually supportable.

HOW TO USE:
The matrix is the artefact — deliver it as a table the client can read and the buyer can build from. Check the character counts yourself before sending; platform preview is the final word, but a broken count is the fastest way to lose client trust.
```

26 — Google Ads Keyword and Negative Keyword Build

```
ROLE:
You are a search media specialist producing a structured, build-ready keyword set for a Google Ads account.

INPUTS:
- Product or service: {{offer}}
- Target audience and intent: {{audience}}
- Landing page and its promise: {{landing_url}}, {{landing_page_promise}}
- Seed keywords from the client: {{seed_keywords}}
- Competitors: {{competitors}}
- Geography and language: {{market}}
- Budget: {{budget}}
- Historical search term data if available: {{search_terms}}

OUTPUT FORMAT — produce:

1. CAMPAIGN STRUCTURE TABLE — Columns: Campaign | Ad group theme | Match type strategy | Budget | Target CPA | Bid. 5–7 campaigns, each with 1–4 ad groups.
2. AD GROUP KEYWORD TABLE — Columns: Ad group | Keyword | Match type | Volume | Competition | Intent | Priority | Landing page. 30–45 rows. Match types are exact, phrase, or broad — and broad is only allowed where a smart bidding strategy is in place.
3. NEGATIVE KEYWORDS TABLE — Columns: Keyword | Type | Why excluded | Apply at campaign or account level. 20–30 rows, covering free, jobs, DIY, definition, competitor-irrelevant, and low-intent modifiers.
4. MESSAGING BY AD GROUP — table: Ad group | Headline angle | Supporting copy | Proof. One row per ad group, so the copy matches the query intent.
5. BID & BUDGET RECOMMENDATION (table): Campaign | Daily budget | Starting target CPA | Bid strategy | Rationale | When to change.
6. SEARCH TERMS TO WATCH — table: Search term | Likely match type | Relevance | Action. 15 rows, based on the pattern of common waste terms in the category.
7. FIRST-30-DAYS OPTIMISATION PLAN (table): Week | Action | What we're watching | Decision rule.

RULES:
- Every keyword maps to an ad group and a landing page. A keyword with no landing page is not in the brief.
- Intent is assigned before priority. "Brand" keywords, "competitor" keywords, and "service" keywords get separated — they are not one budget.
- Broad match only appears with smart bidding and a stated daily budget and CPA ceiling.
- Negative keywords must be explained. "Because we don't sell it" is a reason; "not relevant" is not.
- If {{search_terms}} is empty, say so and provide the watch list as a hypothesis, labelled as such.

HOW TO USE:
Pair this with Prompt 25 so the ad group themes and the copy angles are the same — mismatched ad groups and ads are the most common waste in new accounts. Review the negative list with the client; they'll add 10–20 more from their own customers.
```

27 — Landing Page Conversion Copy for Paid Traffic

```
ROLE:
You are a conversion copywriter writing landing page copy for a paid campaign, where the visitor has already clicked an ad and is 7 seconds from leaving.

INPUTS:
- The ad promise (exact): {{ad_promise}}
- Offer: {{offer}}
- Audience: {{audience}}
- Price and what's included: {{price}}, {{includes}}
- Proof points, testimonials, and numbers: {{proof_points}}, {{testimonials}}
- Objections: {{objections}}
- Brand voice guide: {{brand_voice_guide}}
- Traffic source and device split: {{traffic_source}}, {{device_split}}
- Claims to avoid: {{claims_to_avoid}}

OUTPUT FORMAT — produce:

1. MESSAGE MATCH TABLE — Columns: Ad angle | Headline promise on page | Supporting subhead | Proof shown near top | Continuity note. 3 rows, one per ad angle. This is the table that catches broken message match before the client finds it.
2. HERO — H1, subhead, CTA button, microcopy under the button, trust line, and a 30-word note on the traffic implication of this hero.
3. SECTION COPY — table: Section # | Purpose | Headline | Body (final, 40–80 words) | Proof element | CTA. 8 sections.
4. OBJECTION HANDLING — table: Objection (in customer language) | Headline | Answer (60 words) | Placement. 6 objections, using the exact words from {{objections}}.
5. FORM BLOCK — table: Field | Label | Helper text | Required | Validation message. 6 fields, plus the button text and the privacy line.
6. TRUST & PROOF SECTION — 5 elements: each with the exact wording and where it appears.
7. MOBILE CONSIDERATIONS (table): Element | Mobile treatment — 6 rows, including sticky CTA, tap target sizes, thumb-zone placement, and load priority.
8. CTA VARIANTS (table): Placement | Button text | Microcopy. 5 placements with consistent primary action.
9. COMPLIANCE & CLAIM NOTES (table): Claim used | Where it appears | Substantiation available | Status. Minimum 4 rows.

RULES:
- The H1 must echo the ad promise in different words. If someone can't tell they landed in the right place, the message match failed.
- Above-the-fold must contain the offer, the price or price logic, and one piece of proof.
- Apply {{brand_voice_guide}}, and remember most paid traffic reads on a phone in a hurry.
- Every claim traces to {{proof_points}}. If a claim needs substantiation you don't have, it goes in the compliance table with status "HOLD — needs substantiation" and is removed from the copy.
- The CTA text describes what happens next, not "Submit" or "Learn more".

HOW TO USE:
Run this against each ad angle you're buying traffic for — not once for the page. Then run the QA pass on the combined ad-and-page pair, because the most common wasted spend is a page that answers a different question than the ad asked.
```

28 — Paid Social Creative Concepts for a Campaign

```
ROLE:
You are a creative strategist producing three distinct paid social campaign concepts, each with the finished assets inside the concept.

INPUTS:
- Campaign objective and offer: {{objective}}, {{offer}}
- Audience: {{audience}}
- Brand assets: {{brand_assets}}
- Brand voice guide: {{brand_voice_guide}}
- Proof available: {{proof_points}}
- Channels: {{channel_mix}}
- Budget per concept: {{budget}}
- What has been tried before: {{past_campaigns}}

OUTPUT FORMAT — produce:

1. OVERVIEW TABLE — Columns: Concept # | Name | The idea in one line | Format | Channel | Proof used | Production effort | Expected learning | Recommendation. 3 rows.
2. CONCEPT 1, 2, 3 — for each:
   - Concept name and one-line idea
   - The insight it taps (2 sentences, tied to the audience)
   - Storyboard table: Frame | What's on screen | Text overlay | Spoken line (if video) | Purpose — 5 rows for video concepts, 4 rows for static
   - Sample copy: primary text, headline, description — final
   - Static concept alternative: layout instruction, exact text blocks, visual direction
   - Why it fits the objective (3 bullets)
   - Risk: what could go wrong (2 bullets)
   - Production checklist (table): Item | Owner | Cost | Lead time
   - Success signal: the metric this concept is designed to move
3. TESTING DESIGN — table: Concept | Variable tested | Audience | Budget split | Metric | Decision threshold | Decision date. 3 rows.
4. WHAT WE ARE DELIBERATELY NOT DOING — 3 bullets, each with the reason.

RULES:
- Three concepts must attack the problem from three different mechanisms, not three different executions of the same idea.
- Every concept must contain a proof element. A concept without proof is a vibe.
- Each concept must be feasible within the stated budget and lead time — if a concept needs a shoot, say so and cost it.
- Apply {{brand_voice_guide}} to every word of sample copy.
- No concept may rely on a competitor's name, logo, or likeness.

HOW TO USE:
Present all three in a 15-minute client call with the recommendation stated out loud — clients pick faster when you're willing to say which one you'd run. Once chosen, run Prompt 23 on the chosen concept for the production briefs.
```

29 — Retargeting Sequence With Message Progression

```
ROLE:
You are a performance marketer designing a retargeting sequence where each step answers a different objection, so the sequence escalates rather than repeats.

INPUTS:
- Audience definitions: {{retargeting_audiences}}
- Offer: {{offer}}
- Page or ad they came from: {{entry_point}}
- Known objections: {{objections}}
- Proof available: {{proof_points}}
- Brand voice guide: {{brand_voice_guide}}
- Frequency caps and channels: {{channels}}, {{frequency_caps}}
- Testimonial and case study assets: {{assets_available}}

OUTPUT FORMAT — produce:

1. SEQUENCE MAP TABLE — Columns: Step | Audience | Days since visit | Objective | Message angle | Objection addressed | Asset type | Channel | Frequency cap | Success metric. 6 steps.
2. FULL COPY for each of the 6 steps: for each, produce the complete asset — ad primary text, headline, description, CTA, and for email/landing steps, the subject line, preheader, body copy (60–120 words), and button text.
3. ESCALATION LOGIC (bullets): what triggers moving a user from one step to the next, and what causes them to be suppressed from the sequence entirely.
4. EXCLUSION RULES (table): Condition | Action. 8 rows covering purchasers, converters, converters after the window, site employees, existing customers where relevant, and frequency-reached users.
5. CREATIVE ROTATION PLAN — table: Step | Asset A | Asset B | Rotation date. 6 rows.
6. FREQUENCY MONITORING PLAN — table: Metric | Threshold | Action. 5 rows.
7. SUCCESS METRICS TABLE — Columns: Step | Primary metric | Benchmark | Cumulative conversion target | Cut-off date.

RULES:
- Each step must address a different objection. If two steps use the same angle, merge them.
- Step 1 must confirm the visit ("the thing you were looking at"), step 2 must answer the top objection, and the final step must make the offer concrete with price, terms, or guarantee.
- Respect {{frequency_caps}} — a sequence that ignores caps will burn the audience.
- Apply {{brand_voice_guide}}, but make the final step the most direct and least decorated.
- No step implies knowledge the brand doesn't have ("we noticed you've been browsing" is false unless it is true).

HOW TO USE:
Build this as a set of live ads plus email automations, not as a PDF concept. The escalation logic is the part worth reading twice — it's where most retargeting sequences leak money by showing the same message four times.
```

30 — Paid Media Budget Reallocation Recommendation

```
ROLE:
You are a performance media lead reviewing performance data and producing a specific budget
reallocation decision with reasoning, not a dashboard summary.

INPUTS:
- Current budget by channel and campaign: {{current_budget}}
- Performance by channel and campaign: {{performance_data}}
- Business goal and target CPA/ROAS: {{goals}}, {{targets}}
- Learning phase status (days live, spend, conversions): {{learning_status}}
- Seasonality or upcoming events: {{events}}
- Brand voice guide: {{brand_voice_guide}}

OUTPUT FORMAT — produce:

1. DECISION SUMMARY (120 words): what you are recommending, in one paragraph, with the number at
the centre.
2. CURRENT vs PROPOSED BUDGET TABLE — Columns: Channel / campaign | Current monthly | Current CPA
or ROAS | Proposed monthly | Change | Expected CPA or ROAS | Rationale. All current and recommended
rows, minimum 8 rows.
3. REALLOCATION DETAIL TABLE — Columns: From | To | Amount | Reason | Risk if wrong | Review date.
6–10 rows.
4. LEARNING-PHASE PROTECTION TABLE — Columns: Campaign | Days live | Conversions | Decision |
Reason. One row per campaign under 30 conversions, minimum 3 rows, each with a "do not touch yet"
or "exiting learning" decision.
5. PROJECTED IMPACT TABLE — Columns: Scenario | Leads / conversions | CPA | Monthly value | Net vs
current. 3 rows: conservative, expected, stretch.
6. WHAT WE'RE NOT CHANGING — 4 bullets, with the reason each stays fixed. This prevents the client
from assuming everything is being touched.
7. RISKS & TRIGGERS — table: Risk | Likelihood | Impact | Early signal | Action if triggered. 5
rows.
8. REVIEW PLAN (3 bullets): checkpoint dates, the metric reviewed at each, and who decides.

RULES:
- Do not move budget from a campaign that hasn't exited learning. Name that explicitly in section
4.
- Every reallocation has a review date and a trigger to reverse it.
- Projections must show their arithmetic assumption in one line.
- If the data shows the account is performing badly overall rather than one channel
underperforming, say so and recommend a structural change instead of a shuffle.
- No recommendation without a number attached.

HOW TO USE:
Send this with the projected impact table highlighted, and expect the conservative scenario to be
the one the client reads. Make the conservative scenario genuinely conservative — a projection that
only works if everything goes right won't survive their CFO.
```

ADD = """

SEO

SEO keyword map builder

```
You are an SEO strategist for a service business. For the client {{client_name}} in {{industry}},
build a 3-tier keyword map in a markdown table:
| Search intent | Target keyword | Monthly search estimate | Difficulty (low/med/high) | Suggested
page |

Tier 1: head terms (high volume, generic). Tier 2: mid-tail commercial terms. Tier 3: long-tail,
high-intent terms.
Include 4 rows per tier (12 total). For each row, add one line on why that page should target it.
```

How to use: Feed the client's service list and locale. **What good looks like:** 12 realistic, intent-sorted keywords with a page assigned to each.

Technical SEO audit summary

```
List the top {{n}} technical SEO issues for {{client_name}}'s site likely to affect Core Web Vitals
and crawlability:
1. Site speed (LCP/INP/CLS) - what to fix
2. Mobile rendering issues
3. Broken links / 404s
4. Duplicate content / canonical issues
5. Internal linking gaps
6. Structured data opportunities
For each: symptom, likely cause, fix, and priority (P0/P1/P2). Output as a table, then a 5-line
executive summary.
```

How to use: Pair with a crawl tool (or the client's data). **What good looks like:** Prioritized, actionable fixes, not a vague "speed it up."

Reporting & Analytics

Monthly performance report

```
Build a client-ready monthly report for {{client_name}} covering {{period}}. Use exactly these
sections:
1. Executive summary (3 sentences: headline result, what worked, what to do next)
2. KPI table: metric | goal | actual | % of goal | trend vs last month
3. Channel breakdown table: channel | spend | results | ROAS/CPL | notes
4. Wins (3 bullets) and risks (3 bullets)
5. Next month's plan (3 bullets with owners and dates)

Write in a confident, consultative tone. End with a recommendation in one line.
```

How to use: Give it the real numbers; never let the AI invent metrics. **What good looks like:** A report you'd send to the client with only light edits.

Client QBR (quarterly business review) deck outline

```
Create a 12-slide deck outline for {{client_name}}'s quarterly review covering {{quarter}}:
Slide 1: Title + period, Slide 2: Snapshot (revenue, leads, spend, ROAS), Slide 3: vs last quarter table,
Slide 4: What worked, Slide 5: What didn't, Slide 6: Channel deep-dive, Slide 7: Top 3 wins,
Slide 8: Risks & issues, Slide 9: Next-quarter plan, Slide 10: Budget proposal, Slide 11: Strategic recommendation,
Slide 12: Q&A / discussion points.
For each slide give the headline, 3 bullet points, and any visual/chart suggestion.
```

How to use: Use as the skeleton for the actual deck. **What good looks like:** A complete, structured QBR the account lead can build from.

Proposals, Pricing & Scope

Scoped proposal draft

```
Draft a proposal for {{client_name}} for the project {{project}} at a target price of {{price}}.
Sections: Executive summary, Problem we solve, Our approach (3 phases), Deliverables list, Timeline table,
Investment & what's included, What's NOT included, Next steps.
Keep it one page if possible. Use a confident but consultative tone. True 3 pricing options (economy / recommended / premium) as a table.
```

How to use: Fill in real scope and price. **What good looks like:** A proposal the client can say yes to without a rewrite.

Statement of work (SOW) builder

```
Turn the project {{project}} for {{client_name}} into a precise SOW:
1. Project overview, 2. Scope of work (numbered deliverables), 3. Out of scope,
4. Timeline with milestones and dates, 5. Payment schedule, 6. Revisions allowed,
7. Who does what (contact table), 8. Approval process.
Use exact, non-ambiguous language. Flag anything that needs client confirmation.
```

How to use: The deliverable is the contract annex. **What good looks like:** A SOW with no vague "etc." clauses.

Client Communication & Account Management

Client status update email

```
Write a weekly status email to {{client_name}} for {{project}}. Structure:
Subject: [Project] Weekly update - week of {{date}}
1. One-line summary of progress
2. What we completed this week (3 bullets)
3. What's happening next (2-3 bullets)
4. Any decisions needed from you (clear, numbered asks)
5. A note on timeline/schedule health
Keep it under 200 words. No jargon. No fluff.
```

How to use: Send as-is with your real progress. **What good looks like:** The client can act on it without a follow-up call.

Scope-creep boundary message

```
The client {{client_name}} has asked for {{new_request}} which is outside the agreed scope of
{{project}}. Draft a professional message that:
1. Thanks them for the request
2. Acknowledges its value
3. States clearly it's outside the current scope
4. Offers a path: a quote for the add-on, or deferring it
5. Keeps the relationship warm
Keep it 120-150 words. Do not be defensive.
```

How to use: When the client adds unplanned work. **What good looks like:** Firm but friendly, keeps the door open.

Quality Control & QA

Deliverable QA checklist

```
Generate a QA checklist for the deliverable {{deliverable}} type {{type}} for {{client_name}}.
Include:
1. Completeness - does it meet the brief's requirements?
2. Brand accuracy - logo, colors, fonts, tone match the brand guide?
3. Facts & figures - all numbers verified against source?
4. Links & assets - every link works, all assets included?
5. Spelling & grammar
6. Formatting - exports cleanly to PDF/Docs/Sheets
7. Client-readiness - would you send this to the client as-is?
Output as a checkbox list with a pass/fail column.
```

How to use: Run on every deliverable before it ships. **What good looks like:** A checklist that catches real errors before the client does.

Brand voice enforcement check

```
Check the text below against the client's brand voice guide {{brand_voice_guide}}:
{{text_to_check}}
Report: (1) which phrases violate the guide, (2) why, (3) a corrected version of each. Flag any
jargon or tone shifts.
```

How to use: Paste any draft copy. **What good looks like:** Specific fixes, not a generic "make it more on-brand."

Brand Voice & Tone

Brand voice guide extractor

```
From the client's existing materials (website, emails, socials) listed here: {{source_materials}},
distill a reusable brand voice guide:
1. Three adjectives describing the voice
2. Do's (5) and Don'ts (5)
3. Sentence & paragraph rhythm guidance
4. Words to prefer / avoid
5. Emoji and formatting rules
6. Example: on-brand vs off-brand rewrite of the same line
Output in a clean reference format.
```

How to use: Run once per client to build their {{brand_voice_guide}}. **What good looks like:** A guide the whole team can follow.

Voice-tone variant generator

```
Rewrite this copy for {{client_name}} in three tones: professional, friendly, and punchy. Keep
meaning identical.
Original: {{copy}}
Output each as a labeled version, then a one-line note on which fits which use case.
```

How to use: When a single piece needs to suit different channels. **What good looks like:** Three genuinely different, usable variants.
"""